



<https://algs4.cs.princeton.edu>

2.3 QUICKSORT

- ▶ *quicksort*
- ▶ *selection*
- ▶ *duplicate keys*
- ▶ *system sorts*

Two classic sorting algorithms: mergesort and quicksort

Critical components in the world's computational infrastructure.

- Full scientific understanding of their properties has enabled us to develop them into practical system sorts.
- Quicksort honored as one of top 10 algorithms of 20th century in science and engineering.

Mergesort. [last lecture]



Quicksort. [this lecture]



Quicksort t-shirt

```

k) lo = i + 1; else return a[i]; } return a[lo]; } p
mpareTo(w) < 0); } private static void exch(Object[] a,
private static boolean isSorted(Comparable[] a) { return
ted(Comparable[] a, int lo, int hi) { for (int i = lo + 1;
n true; } private static void show(Comparable[] a) { for (in
public static void main(String[] args) { String[] a = StdIn.re
or (int i = 0; i < a.length; i++) { String ith = (String) Quick.
public class Quick { public static void sort(Comparable[] a) { S1
static void sort(Comparable[] a, int lo, int hi) { if (hi <= lo)
(a, lo, j-1); sort(a, :                     assert isSorted(a, lo, hi);
|o, int hi) { int i = lo:                         + 1; Comparable v = a[i]
|ak; while (less(v, a[-
|ic static Comparable se:                         : lo) break; if (i >= j)
|ected element out of b:                         : le[] a, int k) { if (k
|ition(a, lo, hi); if (i:                         : dRandom.shuffle(a); int
boolean less(Comparable v, Comparable w) { return (v.compare
int j) { Object swap = a[i]; a[i] = a[j]; a[j] = swap; } pr
n isSorted(a, 0, a.length - 1); } private static boolean is
1; i <= hi; i++) if (less(a[i], a[i-1])) return false; re*
|int i = 0; i < a.length; i++) { StdOut.println(a[i]); }
= StdIn.readStrings(); Quick.sort(a); show(a); StdOut
ring) Quick.select(a, i); StdOut.println(ith); } } }
ndom.shuffle(a); sort(a, 0, a.length - 1); } priv
eturn; int j = partition(a, lo, hi); sort(a, lo
static int partition(Cor
) { while (less(a[++i],
|a, i, j); } exch(a, lo,
|th) { throw new Runtime
0, hi = a.length - 1; \n
else return a[i]; } ret
mpareTo(w) < 0); } private static
private static boolean isSorted(
ted(Comparable[] a, int lo, int l
n true; } private static void sh
public static void main(String[]
for (int i = 0; i < a.length; i++)
public class Quick { public static
static void sort(Comparable[] a,
(a, lo, i-1); sort(a, i+1, hi); }

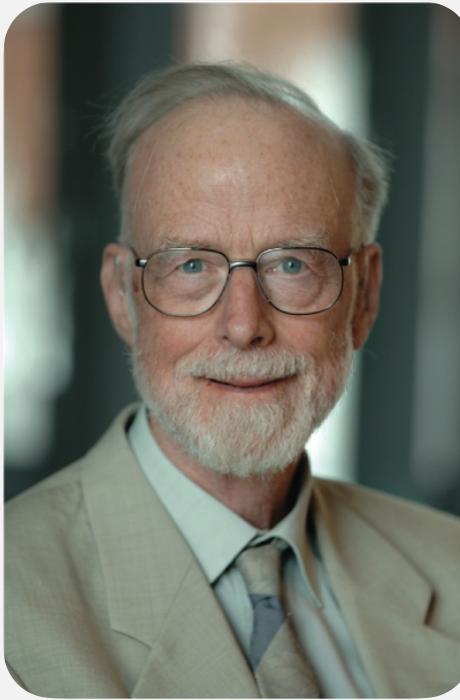
```

CS @ Princeton

A brief history

Tony Hoare.

- Invented quicksort in 1960 to translate Russian into English.
- Learned Algol 60 (and recursion) to implement it.



Tony Hoare
1980 Turing Award

The image contains three panels. The left panel shows the front cover of a book titled 'Algorithms' with a grayscale abstract pattern on the cover. The middle panel shows the front cover of a book titled 'Implementing Quicksort Programs' by Robert Sedgewick, with the subtitle 'Programming Techniques'. The right panel shows a page from the journal 'Acta Informatica' volume 7, pages 327–355 (1977), with the title 'The Analysis of Quicksort Programs*' and author 'Robert Sedgewick'.

ALGORITHM 64
QUICKSORT
C. A. R. HOARE
Elliott Brothers Ltd., Borehamwood, Hertfordshire, Eng.
procedure quicksort (A,M,N); integer M,N;
comment Quicksort is a very fast and convenient method of
sorting an array in the random-access store of a computer. The
entire contents of the store may be sorted, since no extra space is
required. The average number of comparisons made is $2(M-N)$ in
 $(N-M)$, and the average number of exchanges is one sixth this
amount. Suitable refinements of this method will be desirable for
its implementation on an actual computer;
begin integer I,J;
if $M < N$ then begin partition (A,M,N,I,J);
quicksort (A,M,I);
quicksort (A, I, N)
end
quicksort

Programming Techniques S. L. Graham, R. L. Rivest Editors
Implementing Quicksort Programs Robert Sedgewick Brown University

Acta Informatica 7, 327—355 (1977)
© by Springer-Verlag 1977

The Analysis of Quicksort Programs*
Robert Sedgewick
Received January 19, 1976

Summary. The Quicksort sorting algorithm and its best variants are presented and analyzed. Results are derived which make it possible to obtain exact formulas describing the total expected running time of particular implementations on real computers of Quicksort and an improvement called the median-of-three modification. Detailed analysis of the effect of an implementation technique called loop unwrapping is presented. The paper is intended not only to present results of direct practical utility, but also to illustrate the intriguing mathematics which arises in the complete analysis of this important algorithm.

Bob Sedgewick.

- Refined and popularized quicksort in 1970s.
- Analyzed many versions of quicksort.



Bob Sedgewick



2.3 QUICKSORT

- ▶ *quicksort*
- ▶ *selection*
- ▶ *duplicate keys*
- ▶ *system sorts*



Quicksort overview

Step 1. Shuffle the array.

Step 2. Partition the array so that, for some index j :

- Entry $a[j]$ is in place. ← “pivot” or “partitioning item”
- No larger entry to the left of j.
- No smaller entry to the right of j.

Step 3. Sort each subarray recursively.

input	Q	U	I	C	K	S	O	R	T	E	X	A	M	P	L	E
shuffle	K	R	A	T	E	L	E	P	U	I	M	Q	C	X	O	S
partition	E	C	A	I	E	K	L	P	U	T	M	Q	R	X	O	S
	<i>not greater</i>					<i>partitioning item</i>										
sort left	A	C	E	E	I	K	L	P	U	T	M	Q	R	X	O	S
sort right	A	C	E	E	I	K	L	M	O	P	Q	R	S	T	U	X
result	A	C	E	E	I	K	L	M	O	P	Q	R	S	T	U	X



Quicksort partitioning demo

Repeat until i and j pointers cross:

- Scan i from left to right so long as $(a[i] < a[lo])$.
- Scan j from right to left so long as $(a[j] > a[lo])$.
- Exchange $a[i]$ with $a[j]$.



stop i scan because $a[i] \geq a[lo]$

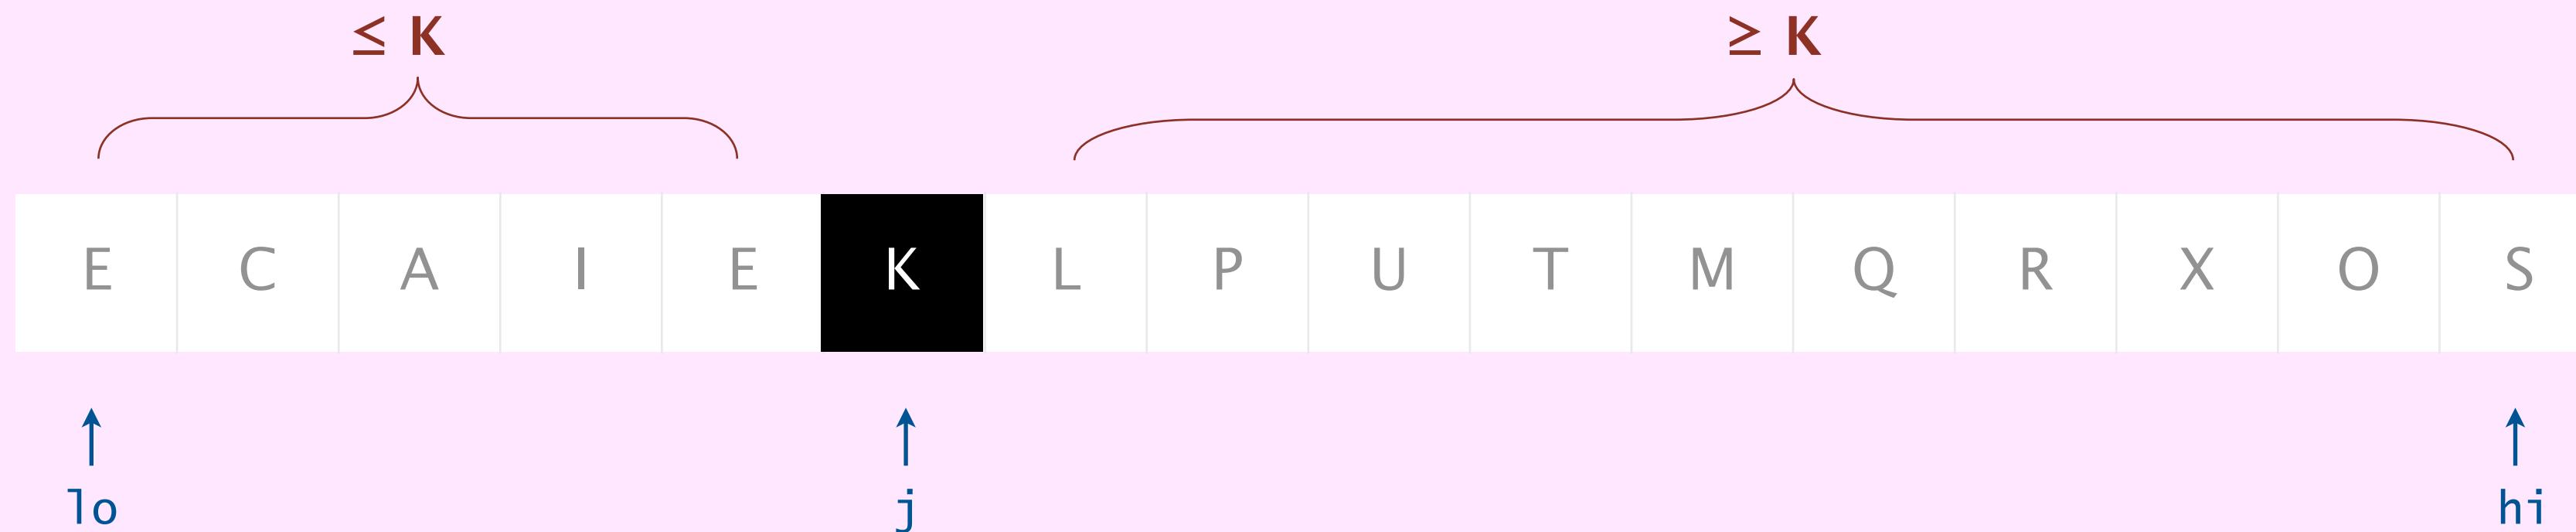


Quicksort partitioning demo

Repeat until i and j pointers cross:

- Scan i from left to right so long as $(a[i] < a[lo])$.
- Scan j from right to left so long as $(a[j] > a[lo])$.
- Exchange $a[i]$ with $a[j]$.

When pointers cross. Exchange $a[lo]$ with $a[j]$.



partitioned!

The music of quicksort partitioning (by Brad Lyon)



New

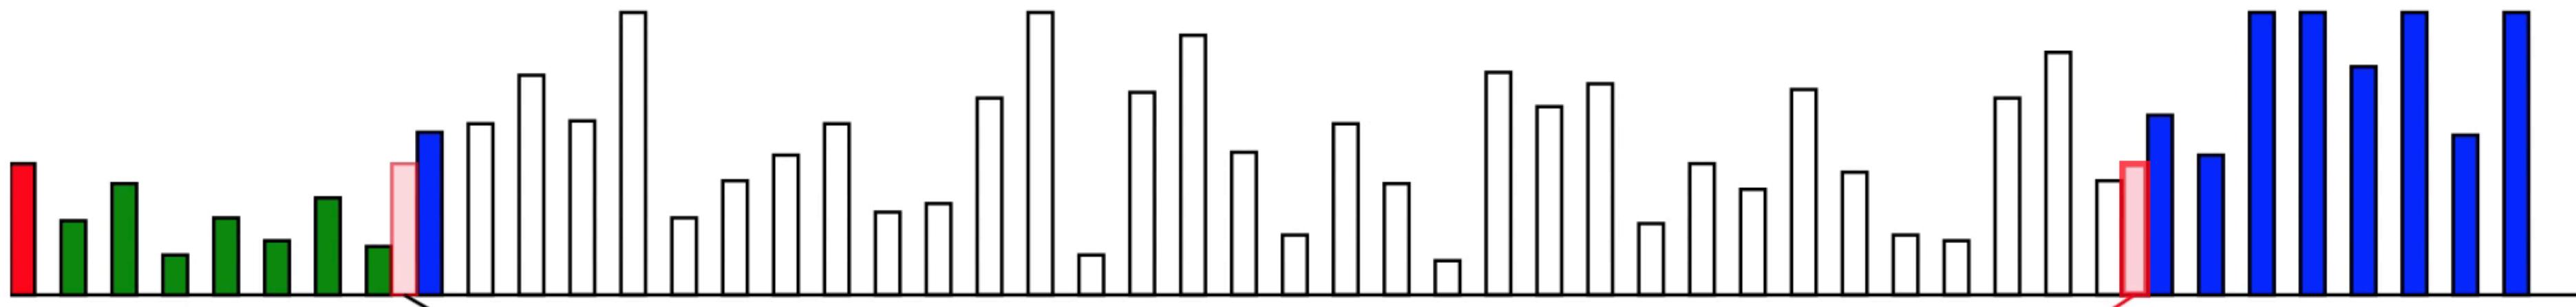
New (Small)

Increasing

Decreasing

Next Step

Do Auto



The value was larger than the pivot, so the lower one waits while the upper one comes down

We will now start coming down from the right

https://learnforeverlearn.com/pivot_music

Quicksort partitioning: Java implementation

```
private static int partition(Comparable[] a, int lo, int hi)
{
    int i = lo, j = hi+1;
    while (true)
    {
        while (less(a[++i], a[lo]))
            if (i == hi) break;

        while (less(a[lo], a[--j]))
            if (j == lo) break;

        if (i >= j) break;
        exch(a, i, j);
    }

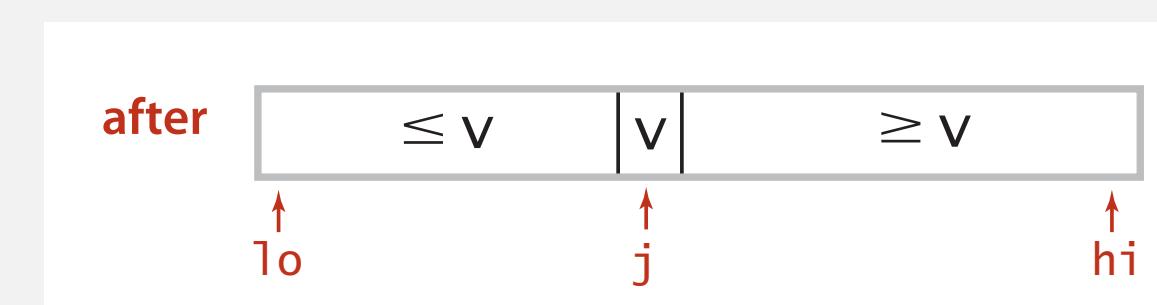
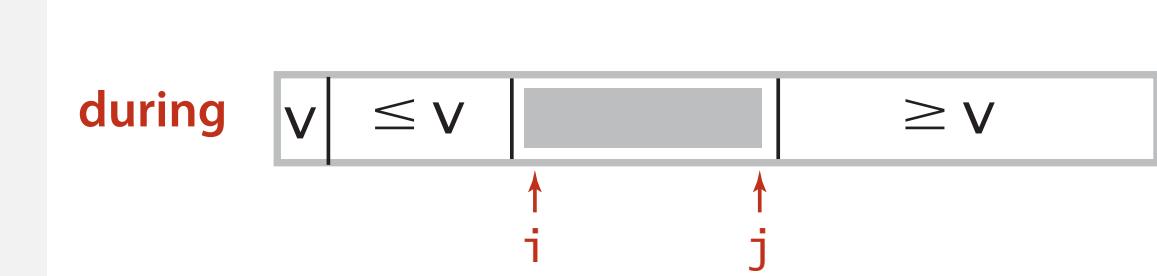
    exch(a, lo, j);           swap with pivot
    return j;                 return index of item now known to be in place
}
```

find item on left to swap

find item on right to swap

check if pointers cross
swap

swap with pivot
return index of item now known to be in place



<https://algs4.cs.princeton.edu/23quick/Quick.java.html>



Quicksort quiz 2

In the worst case, how many compares and exchanges does partition() make to partition a subarray of length n ?

- A. $\sim \frac{1}{2} n$ and $\sim \frac{1}{2} n$
- B. $\sim \frac{1}{2} n$ and $\sim n$
- C. $\sim n$ and $\sim \frac{1}{2} n$
- D. $\sim n$ and $\sim n$

M	A	B	C	D	E	V	W	X	Y	Z
0	1	2	3	4	5	6	7	8	9	10

Quicksort: Java implementation

```
public class Quick
{
    private static int partition(Comparable[] a, int lo, int hi)
    { /* see previous slide */ }

    public static void sort(Comparable[] a)
    {
        StdRandom.shuffle(a);           ← shuffle needed for
        sort(a, 0, a.length - 1);       performance guarantee
                                         (stay tuned)
    }

    private static void sort(Comparable[] a, int lo, int hi)
    {
        if (hi <= lo) return;
        int j = partition(a, lo, hi);
        sort(a, lo, j-1);
        sort(a, j+1, hi);
    }
}
```

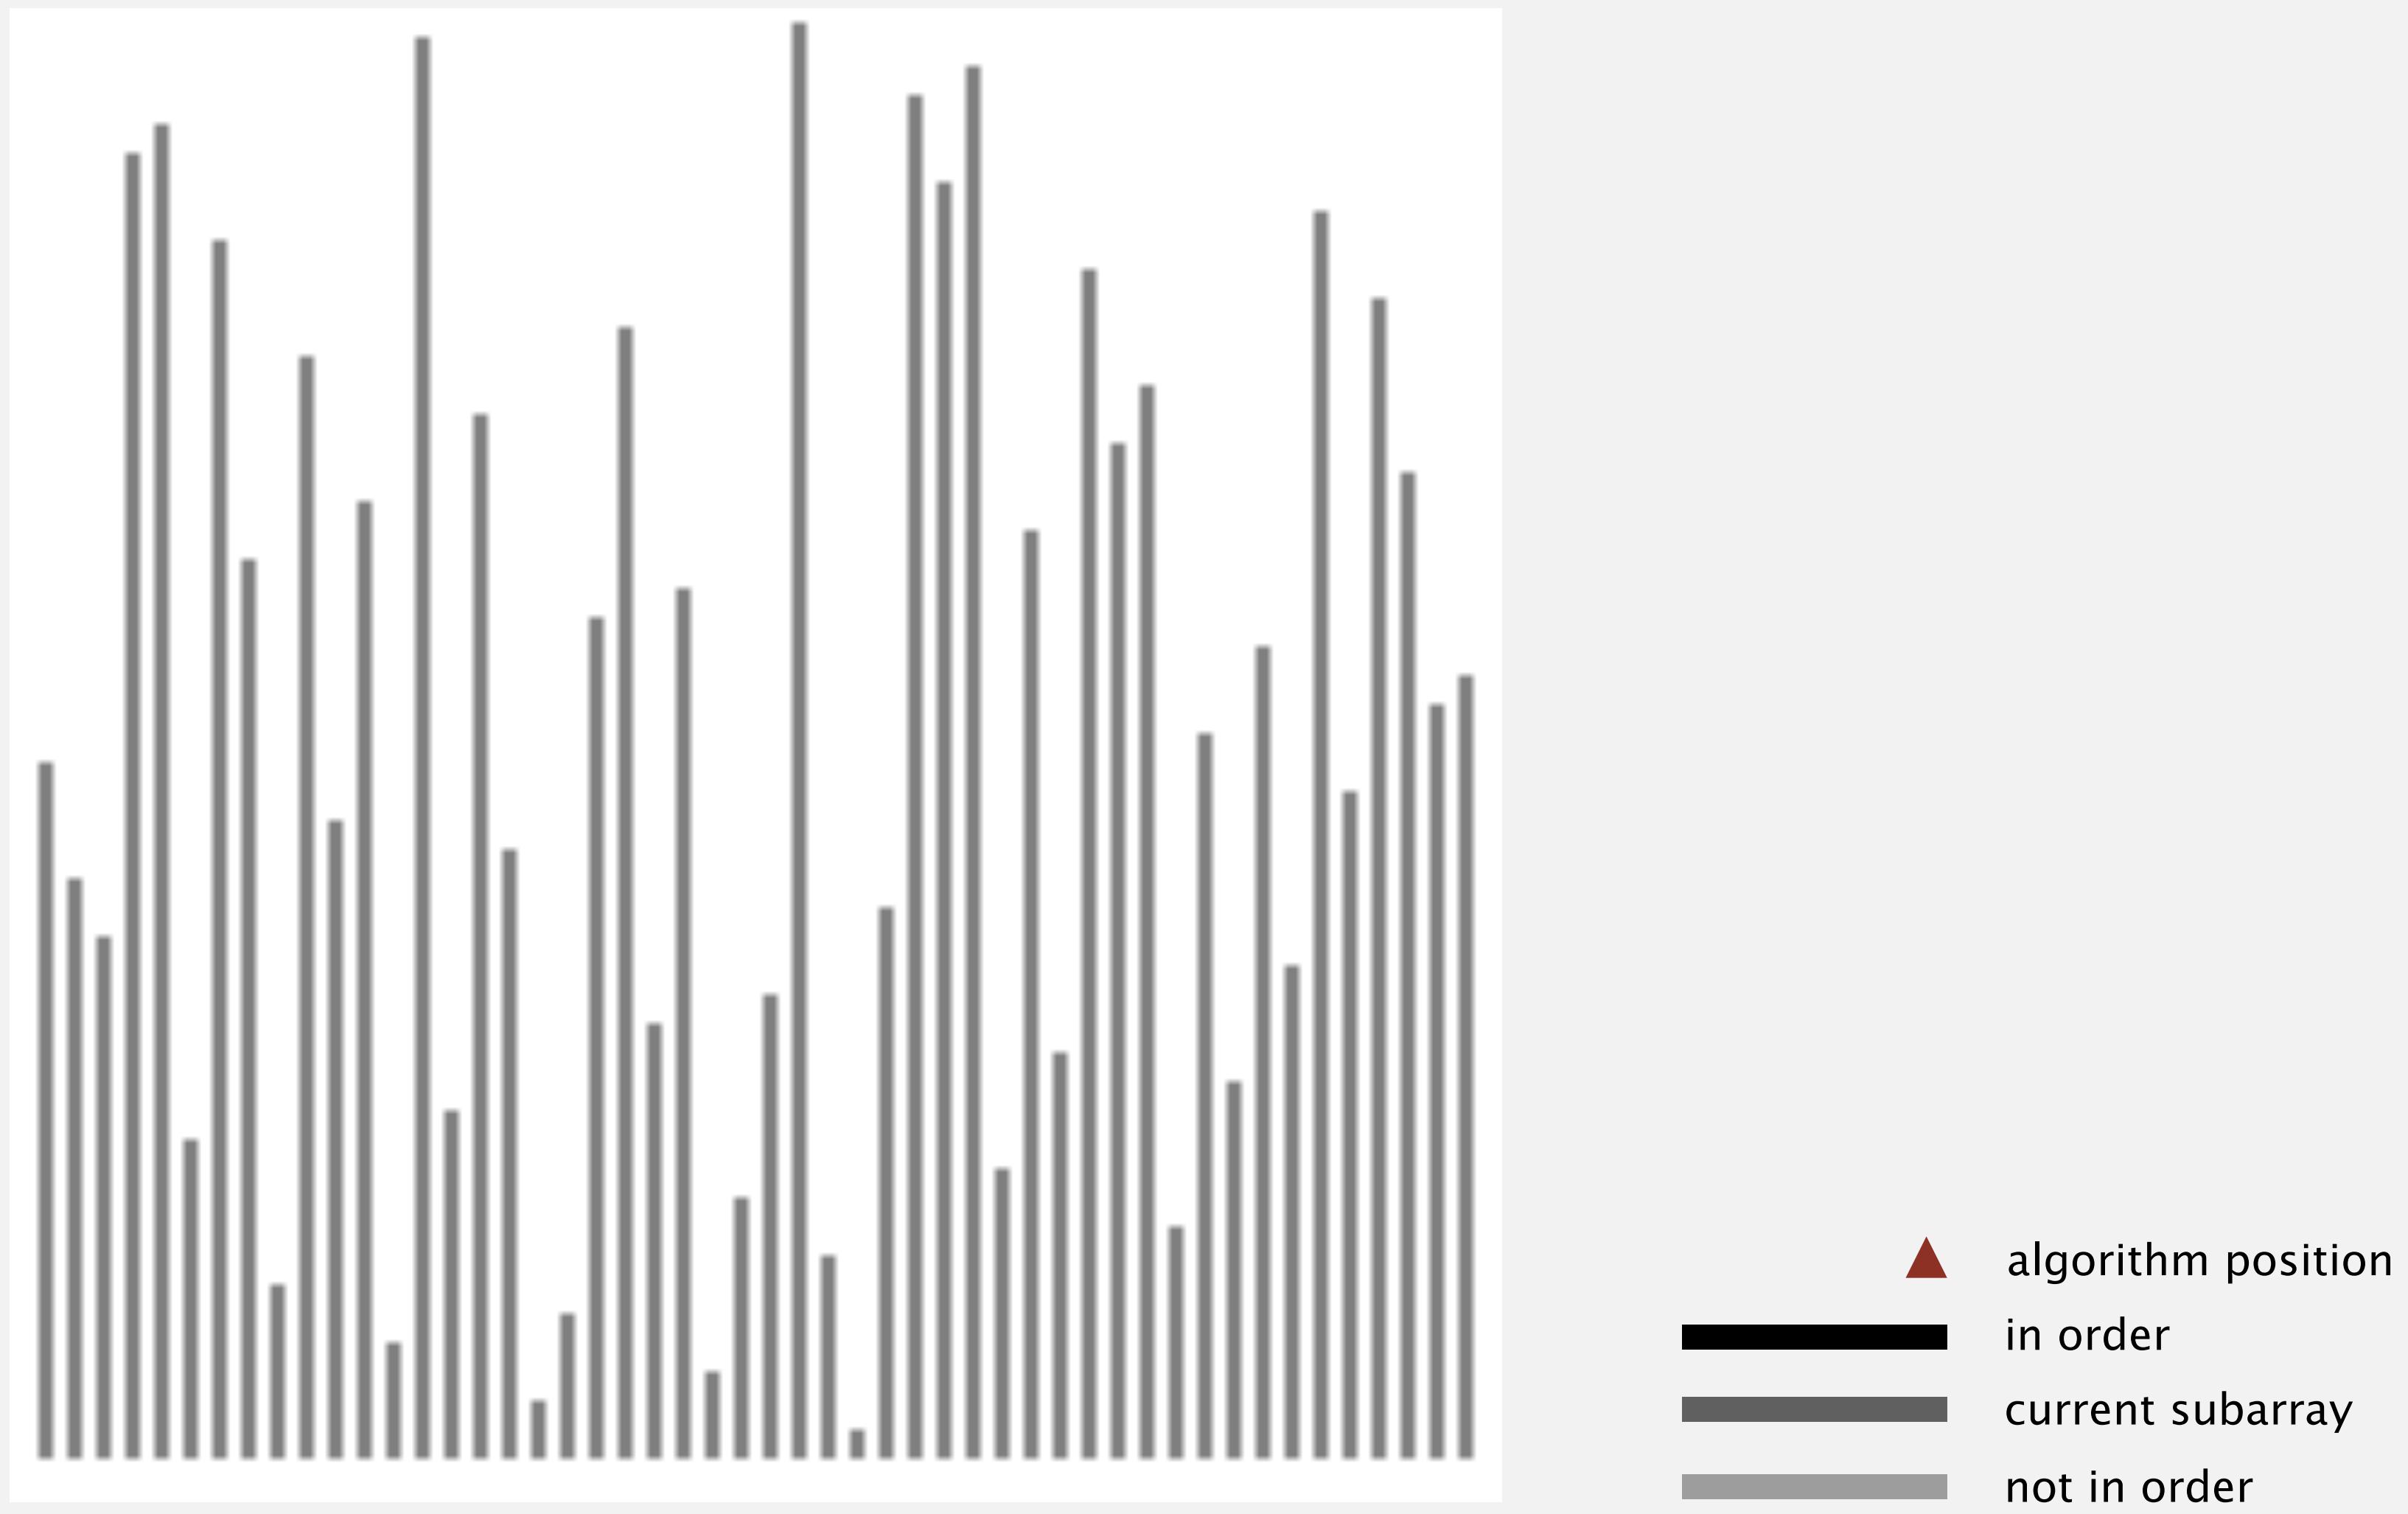
<https://algs4.cs.princeton.edu/23quick/Quick.java.html>

Quicksort trace

<i>initial values</i>	<i>lo</i>	<i>j</i>	<i>hi</i>	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
<i>random shuffle</i>				Q	U	I	C	K	S	O	R	T	E	X	A	M	P	L	E
	0	5	15	E	C	A	I	E	K	L	P	U	I	M	Q	C	X	O	S
	0	3	4	E	C	A	E	I	K	L	P	U	T	M	Q	R	X	O	S
	0	2	2	A	C	E	E	I	K	L	P	U	T	M	Q	R	X	O	S
	0	0	1	A	C	E	E	I	K	L	P	U	T	M	Q	R	X	O	S
	1			A	C	E	E	I	K	L	P	U	T	M	Q	R	X	O	S
	4			A	C	E	E	I	K	L	P	U	T	M	Q	R	X	O	S
	6	6	15	A	C	E	E	I	K	L	P	U	T	M	Q	R	X	O	S
<i>no partition for subarrays of size 1</i>	7	9	15	A	C	E	E	I	K	L	M	O	P	T	Q	R	X	U	S
	7	7	8	A	C	E	E	I	K	L	M	O	P	T	Q	R	X	U	S
	8			A	C	E	E	I	K	L	M	O	P	T	Q	R	X	U	S
	10	13	15	A	C	E	E	I	K	L	M	O	P	S	Q	R	T	U	X
	10	12	12	A	C	E	E	I	K	L	M	O	P	R	Q	S	T	U	X
	10	11	11	A	C	E	E	I	K	L	M	O	P	Q	R	S	T	U	X
	10	10	10	A	C	E	E	I	K	L	M	O	P	Q	R	S	T	U	X
	14	14	15	A	C	E	E	I	K	L	M	O	P	Q	R	S	T	U	X
	15			A	C	E	E	I	K	L	M	O	P	Q	R	S	T	U	X
<i>result</i>				A	C	E	E	I	K	L	M	O	P	Q	R	S	T	U	X
Quicksort trace (array contents after each partition)																			

Quicksort animation

50 random items



<http://www.sorting-algorithms.com/quick-sort>

Quicksort: implementation details

Partitioning in-place. Using an extra array makes partitioning easier (and stable), but it is not worth the cost.

Loop termination. Terminating the loop is more subtle than it appears.

Equal keys. Handling duplicate keys is trickier than it appears. [stay tuned]

Preserving randomness. Shuffling is needed for performance guarantee.

Equivalent alternative. Pick a random pivot in each subarray.



Quicksort: empirical analysis

Running time estimates:

- Home PC executes 10^8 compares/second.
- Supercomputer executes 10^{12} compares/second.

	insertion sort (n^2)			mergesort ($n \log n$)			quicksort ($n \log n$)		
computer	thousand	million	billion	thousand	million	billion	thousand	million	billion
home	instant	2.8 hours	317 years	instant	1 second	18 min	instant	0.6 sec	12 min
super	instant	1 second	1 week	instant	instant	instant	instant	instant	instant

Lesson 1. Good algorithms are better than supercomputers.

Lesson 2. Great algorithms are better than good ones.



Quicksort quiz 3

Why is quicksort typically faster than mergesort in practice?

- A. Fewer compares.
- B. Fewer array acceses.
- C. Both A and B.
- D. Neither A nor B.

Quicksort: worst-case analysis

Worst case. Number of compares is $\sim \frac{1}{2} n^2$.

			a[]														
lo	j	hi	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
			A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
			A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
0	0	14	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
1	1	14	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
2	2	14	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
3	3	14	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
4	4	14	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
5	5	14	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
6	6	14	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
7	7	14	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
8	8	14	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
9	9	14	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
10	10	14	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
11	11	14	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
12	12	14	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
13	13	14	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
14		14	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
			A	B	C	D	E	F	G	H	I	J	K	L	M	N	O

after random shuffle

Quicksort: worst-case analysis

Worst case. Number of compares is $\sim \frac{1}{2} n^2$.

lo	j	hi	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
A	B	C	D	E	F	G	H	I	J	K	L	M	N	O			
A	B	C	D	E	F	G	H	I	J	K	L	M	N	O			

← after random shuffle

Good news. Worst case for quicksort is mostly irrelevant in practice.

- Exponentially small chance of occurring.
(unless bug in shuffling or no shuffling)
- More likely that computer is struck by lightning bolt during execution.



Quicksort: probabilistic analysis

Proposition. The expected number of compares C_n to quicksort an array of n distinct keys is $\sim 2n \ln n$ (and the number of exchanges is $\sim \frac{1}{3} n \ln n$).

Recall. Any algorithm with the following structure takes $\Theta(n \log n)$ time.

```
public static void f(int n)
{
    if (n == 0) return;
    f(n/2);           ← solve two problems
    f(n/2);           ← of half the size
    linear(n);        ←————— do  $\Theta(n)$  work
}
```

Intuition. Each partitioning step divides the problem into two subproblems, each of approximately one-half the size.

↑
probabilistically “close enough”

Quicksort: probabilistic analysis

Proposition. The expected number of compares C_n to quicksort an array of n distinct keys is $\sim 2n \ln n$ (and the number of exchanges is $\sim \frac{1}{3} n \ln n$).

Pf. C_n satisfies the recurrence $C_0 = C_1 = 0$ and for $n \geq 2$:

$$C_n = \underset{\text{partitioning}}{\downarrow} (n+1) + \left(\frac{C_0 + C_{n-1}}{n} \right) + \left(\frac{C_1 + C_{n-2}}{n} \right) + \dots + \left(\frac{C_{n-1} + C_0}{n} \right)$$

left right
 ↓ ↓
 n n

- Multiply both sides by n and collect terms: partitioning probability

$$nC_n = n(n+1) + 2(C_0 + C_1 + \dots + C_{n-1})$$

- Subtract from this equation the same equation for $n - 1$:

$$n C_n - (n-1) C_{n-1} = 2n + 2 C_{n-1}$$

- Rearrange terms and divide by $n(n + 1)$:

$$\frac{C_n}{n+1} = \frac{C_{n-1}}{n} + \frac{2}{n+1}$$

analysis beyond
scope of this course

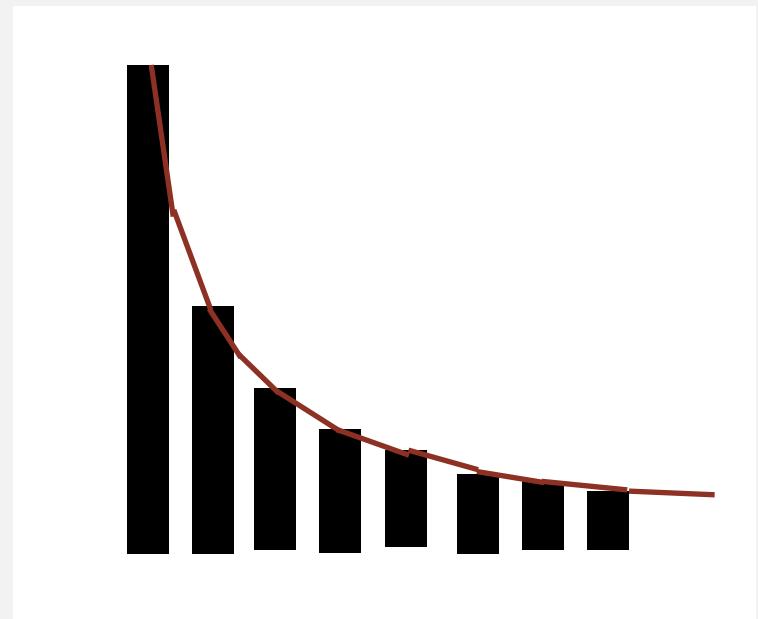
Quicksort: probabilistic analysis

- Repeatedly apply previous equation:

$$\begin{aligned}\frac{C_n}{n+1} &= \frac{C_{n-1}}{n} + \frac{2}{n+1} \\ &= \frac{C_{n-2}}{n-1} + \frac{2}{n} + \frac{2}{n+1} \quad \leftarrow \text{substitute previous equation} \\ &= \frac{C_{n-3}}{n-2} + \frac{2}{n-1} + \frac{2}{n} + \frac{2}{n+1} \\ &= \frac{2}{3} + \frac{2}{4} + \frac{2}{5} + \dots + \frac{2}{n+1}\end{aligned}$$

- Approximate sum by an integral:

$$\begin{aligned}C_n &= 2(n+1) \left(\frac{1}{3} + \frac{1}{4} + \frac{1}{5} + \dots + \frac{1}{n+1} \right) \\ &\sim 2(n+1) \int_3^{n+1} \frac{1}{x} dx\end{aligned}$$



- Finally, the desired result:

$$C_n \sim 2(n+1) \ln n \approx 1.39 n \lg n$$

Quicksort properties

Quicksort analysis summary.

- Expected number of compares is $\sim 1.39 n \log_2 n$.
[standard deviation is $\sim 0.65 n$]
- Expected number of exchanges is $\sim 0.23 n \log_2 n$. ← fewer array accesses than mergesort
- Min number of compares is $\sim n \log_2 n$. ← never fewer than mergesort
- Max number of compares is $\sim \frac{1}{2} n^2$. ← but never happens

Context. Quicksort is a (Las Vegas) randomized algorithm.

- Guaranteed to be correct.
- Running time depends on outcomes of random coin flips (shuffle).



Quicksort properties

Proposition. Quicksort is an **in-place** sorting algorithm.

- Partitioning: $\Theta(1)$ extra space.
- Function-call stack: $\Theta(\log n)$ extra space (with high probability).

can guarantee $\Theta(\log n)$ depth by recurring
on smaller subarray before larger subarray
(but this requires using an explicit stack)

Proposition. Quicksort is **not stable**.

Pf. [by counterexample]

i	j	0	1	2	3
		B_1	C_1	C_2	A_1
1	3	B_1	C_1	C_2	A_1
1	3	B_1	A_1	C_2	C_1
0	1	A_1	B_1	C_2	C_1

Quicksort: practical improvements

Insertion sort small subarrays.

- Even quicksort has too much overhead for tiny subarrays.
- Cutoff to insertion sort for ≈ 10 items.

```
private static void sort(Comparable[] a, int lo, int hi)
{
    if (hi <= lo + CUTOFF - 1)
    {
        Insertion.sort(a, lo, hi);
        return;
    }

    int j = partition(a, lo, hi);
    sort(a, lo, j-1);
    sort(a, j+1, hi);
}
```

Quicksort: practical improvements

Median of sample.

- Best choice of pivot item = median.
- Estimate true median by taking median of sample.
- Median-of-3 (random) items.



$\sim 12/7 n \ln n$ compares (14% fewer)
 $\sim 12/35 n \ln n$ exchanges (3% more)

```
private static void sort(Comparable[] a, int lo, int hi)
{
    if (hi <= lo) return;

    int median = medianOf3(a, lo, lo + (hi - lo)/2, hi);
    swap(a, lo, median);

    int j = partition(a, lo, hi);
    sort(a, lo, j-1);
    sort(a, j+1, hi);
}
```



2.3 QUICKSORT

- ▶ *quicksort*
- ▶ ***selection***
- ▶ *duplicate keys*
- ▶ *system sorts*

Selection

Goal. Given an array of n items, find item of rank k .

Ex. Min ($k = 0$), max ($k = n - 1$), median ($k = n/2$).

Applications.

- Order statistics.
- Find the “top k .”

Use complexity theory as a guide.

- Easy $O(n \log n)$ algorithm. How?
- Easy $O(n)$ algorithm for $k = 0, 1, 2$. How?
- Easy $\Omega(n)$ lower bound. Why?

Which is true?

- $O(n)$ algorithm? [is there a linear-time algorithm?]
- $\Omega(n \log n)$ lower bound? [is selection as hard as sorting?]

Quickselect demo



Partition array so that for some j :

- Entry $a[j]$ is in place.
- No larger entry to the left of j .
- No smaller entry to the right of j .

Repeat in **one** subarray, depending on j ; stop when j equals k .

select element of rank $k = 5$

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
50	21	28	65	39	59	56	22	95	12	90	53	32	77	33

$k = 5$

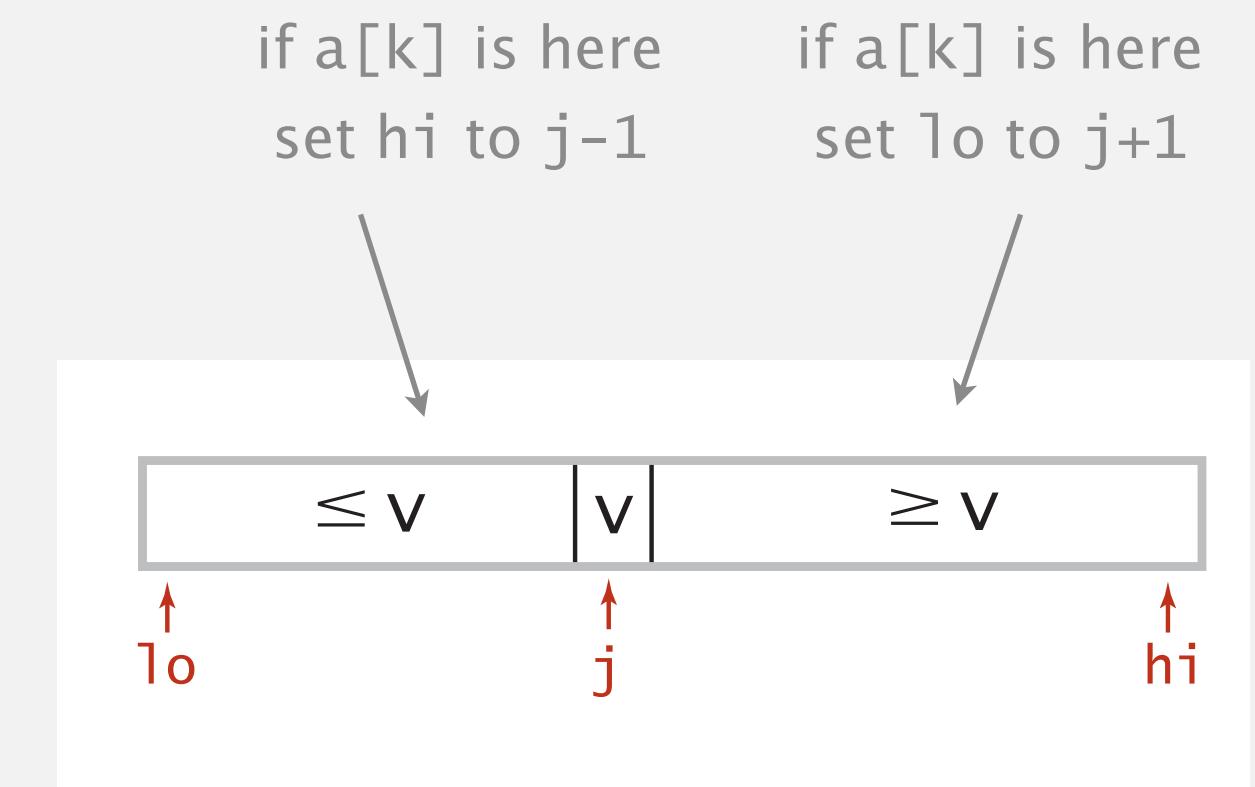
Quickselect

Partition array so that for some j :

- Entry $a[j]$ is in place.
- No larger entry to the left of j .
- No smaller entry to the right of j .

Repeat in **one** subarray, depending on j ; stop when j equals k .

```
public static Comparable select(Comparable[] a, int k)
{
    StdRandom.shuffle(a);
    int lo = 0, hi = a.length - 1;
    while (hi > lo)
    {
        int j = partition(a, lo, hi);
        if (j < k) lo = j + 1;
        else if (j > k) hi = j - 1;
        else return a[k];
    }
    return a[k];
}
```



Quickselect: probabilistic analysis

Proposition. The expected number of compares C_n to quickselect the item of rank k in an array of length n is $\Theta(n)$.

Intuition. Each partitioning step approximately halves the length of the array.

Recall. Any algorithm with the following structure takes $\Theta(n)$ time.

```
public static void f(int n)
{
    if (n == 0) return;
    linear(n);      ← do  $\Theta(n)$  work
    f(n/2);        ← solve one subproblem of half the size
}
```

$$n + n/2 + n/4 + \dots + 1 \sim 2n$$

Careful analysis yields: $C_n \sim 2n + 2k \ln(n/k) + 2(n-k) \ln(n/(n-k))$

$$\begin{aligned} &\leq (2 + 2 \ln 2) n \\ &\approx 3.38 n \end{aligned}$$

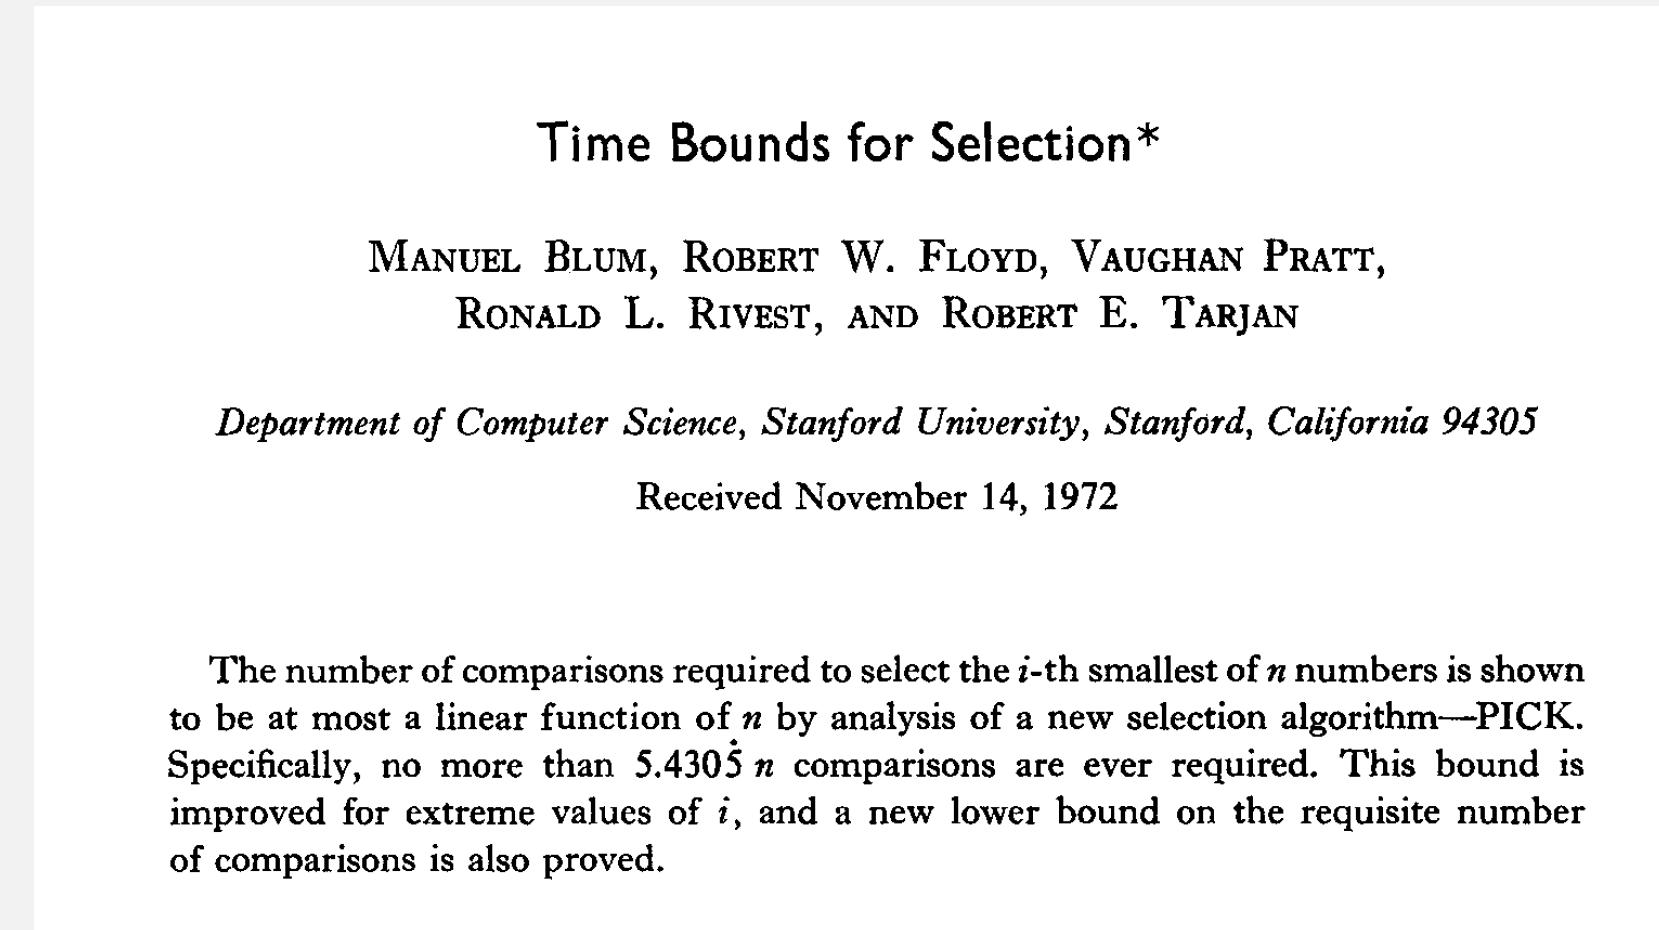
probabilistically “close enough”

max occurs for median ($k = n/2$)

Theoretical context for selection

Q. Compare-based selection algorithm that makes $\Theta(n)$ compares in the **worst case**?

A. Yes! [ingenious divide-and-conquer]



$$T(n) = T(n/5) + T(7n/10) + \Theta(n)$$

↑
find pivot ↑
that eliminates
30% of items

Caveat. Constants are high \Rightarrow not used in practice.

Use theory as a guide.

- Open problem: **practical** algorithm that makes $\Theta(n)$ compares in the worst case.
- Until one is discovered, use quickselect (if you don't need a full sort).



2.3 QUICKSORT

- ▶ *quicksort*
- ▶ *selection*
- ▶ *duplicate keys*
- ▶ *system sorts*

Duplicate keys

Often, purpose of sort is to bring items with equal keys together.

- Sort population by age.
- Remove duplicates from mailing list.
- Sort job applicants by college attended.

Typical characteristics of such applications.

- Huge array.
- Small number of key values.

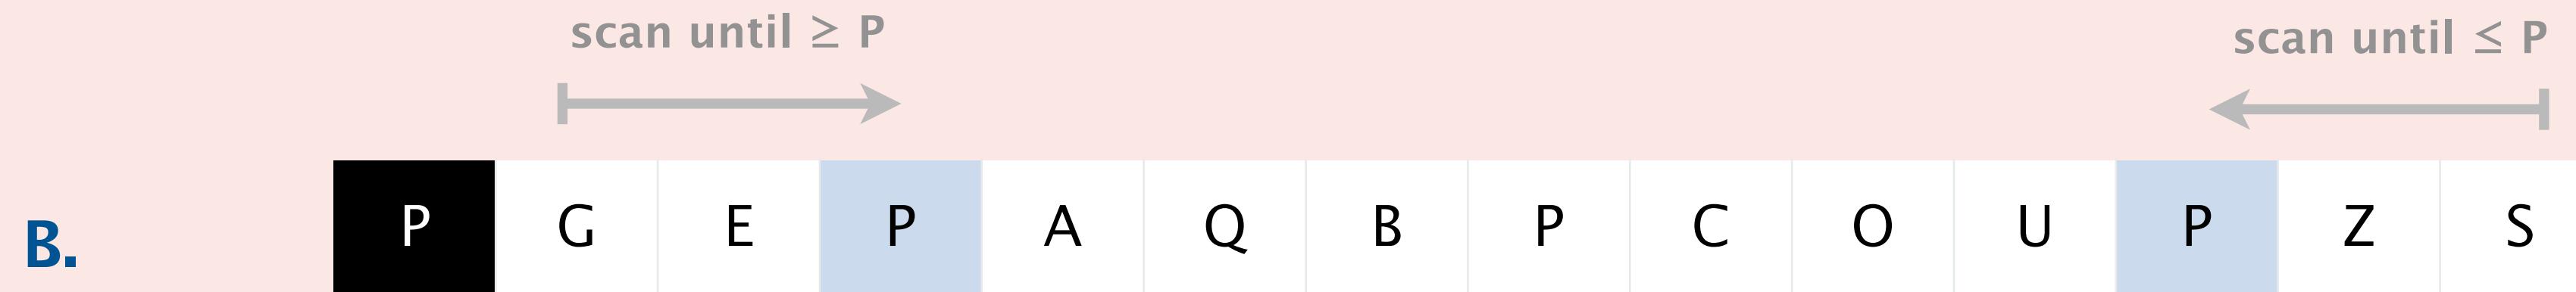
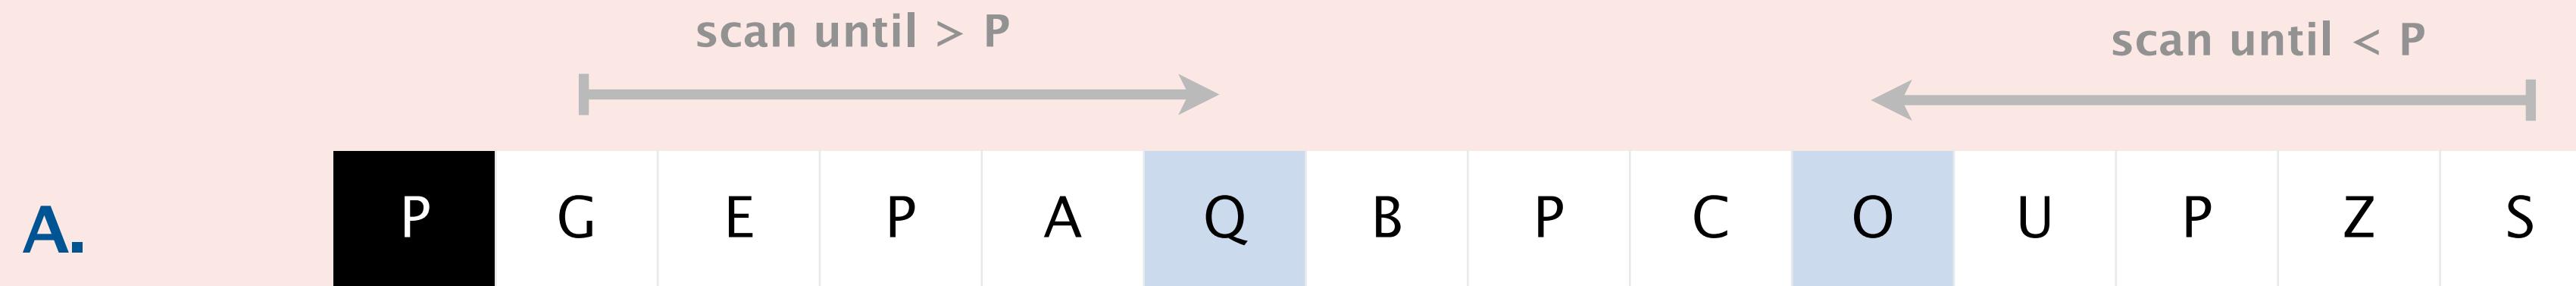
Chicago	09:25:52
Chicago	09:03:13
Chicago	09:21:05
Chicago	09:19:46
Chicago	09:19:32
Chicago	09:00:00
Chicago	09:35:21
Chicago	09:00:59
Houston	09:01:10
Houston	09:00:13
Phoenix	09:37:44
Phoenix	09:00:03
Phoenix	09:14:25
Seattle	09:10:25
Seattle	09:36:14
Seattle	09:22:43
Seattle	09:10:11
Seattle	09:22:54

↑
key



Quicksort quiz 4

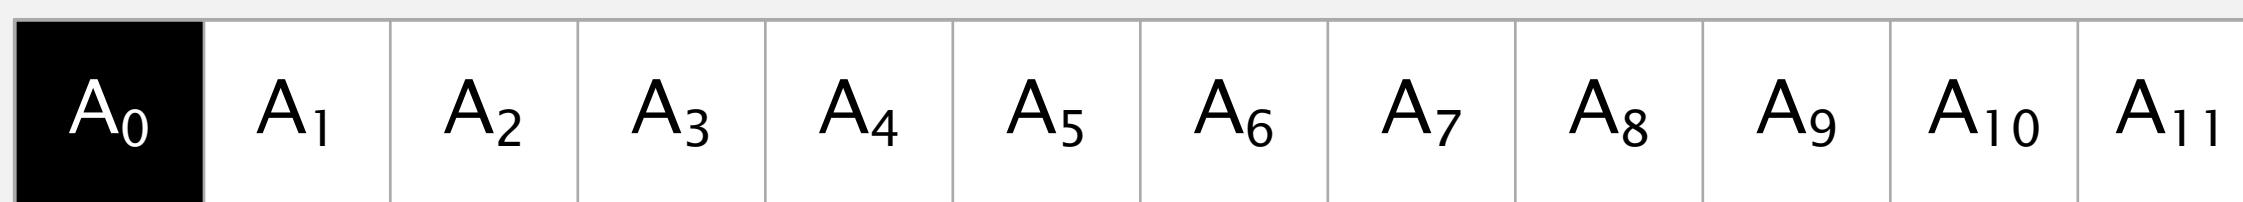
When partitioning, how to handle keys equal to pivot?



C. Either A or B.

War story (system sort in C)

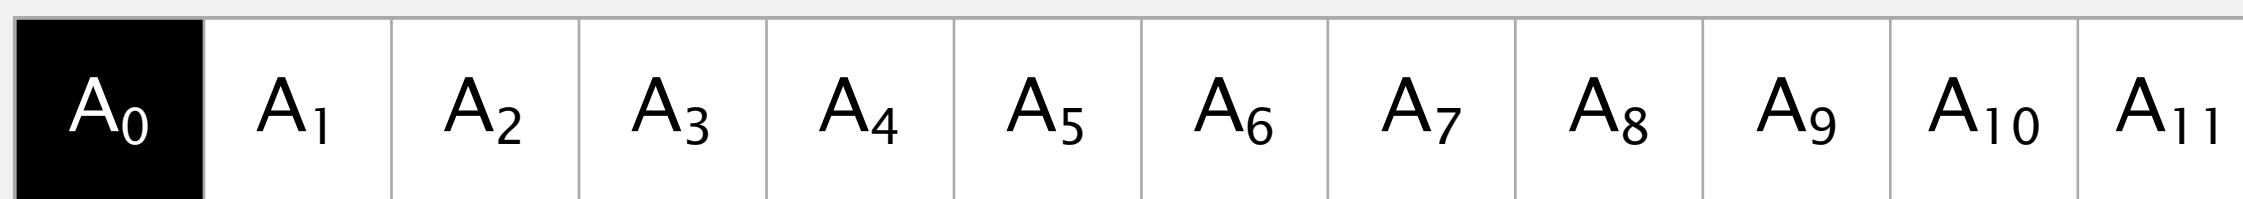
Bug. A qsort() call in C that should have taken seconds was taking minutes to sort a random array of 0s and 1s.



skip over equal keys

↑
i

↑
j



stop scan on equal keys

↑
i

↑
j

Duplicate keys: partitioning strategies

Bad. Don't stop scans on equal keys.

[$\Theta(n^2)$ compares when all keys equal]

B A A B A B B **B** C C C

A A A A A A A A A **A**

Good. Stop scans on equal keys.

[$\sim n \log_2 n$ compares when all keys equal]

B A A B A **B** C C B C B

A A A A A **A** A A A A A

Better. Put all equal keys in place. How?

[$\sim n$ compares when all keys equal]

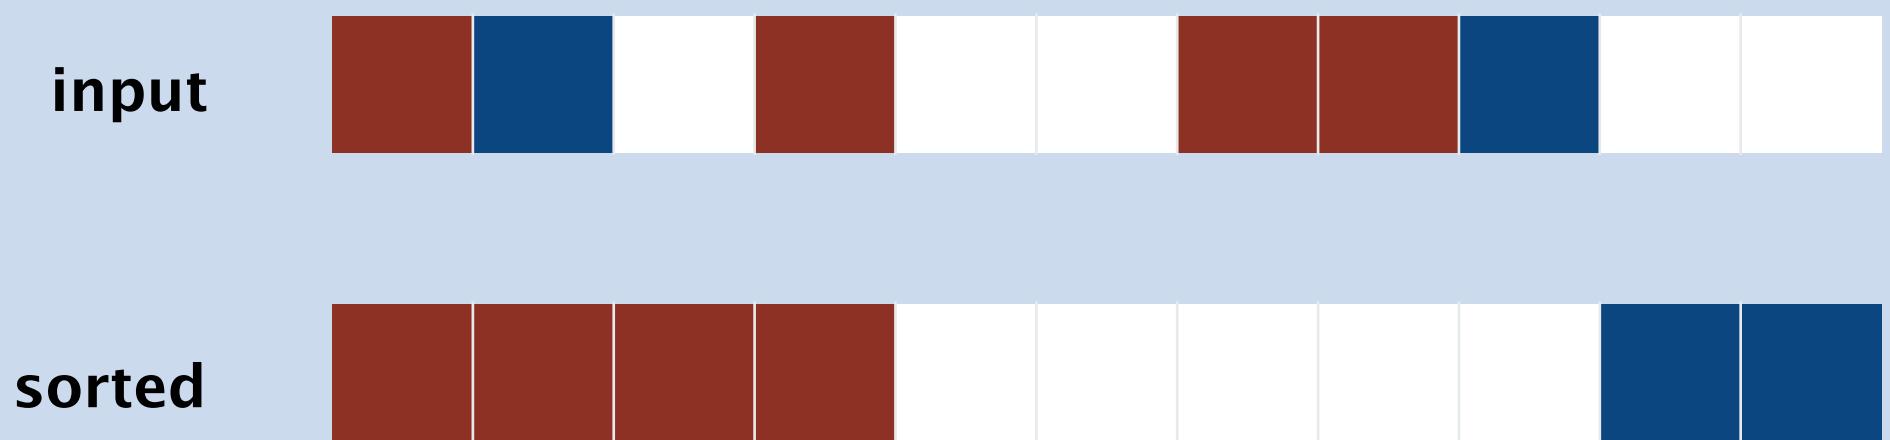
A A A **B** B B B B C C C

A A A A A A A A A A A

DUTCH NATIONAL FLAG PROBLEM



Problem. [Edsger Dijkstra] Given an array of n buckets, each containing a red, white, or blue pebble, sort them by color.



Operations allowed.

- $swap(i, j)$: swap the pebble in bucket i with the pebble in bucket j .
- $getColor(i)$: determine the color of the pebble in bucket i .

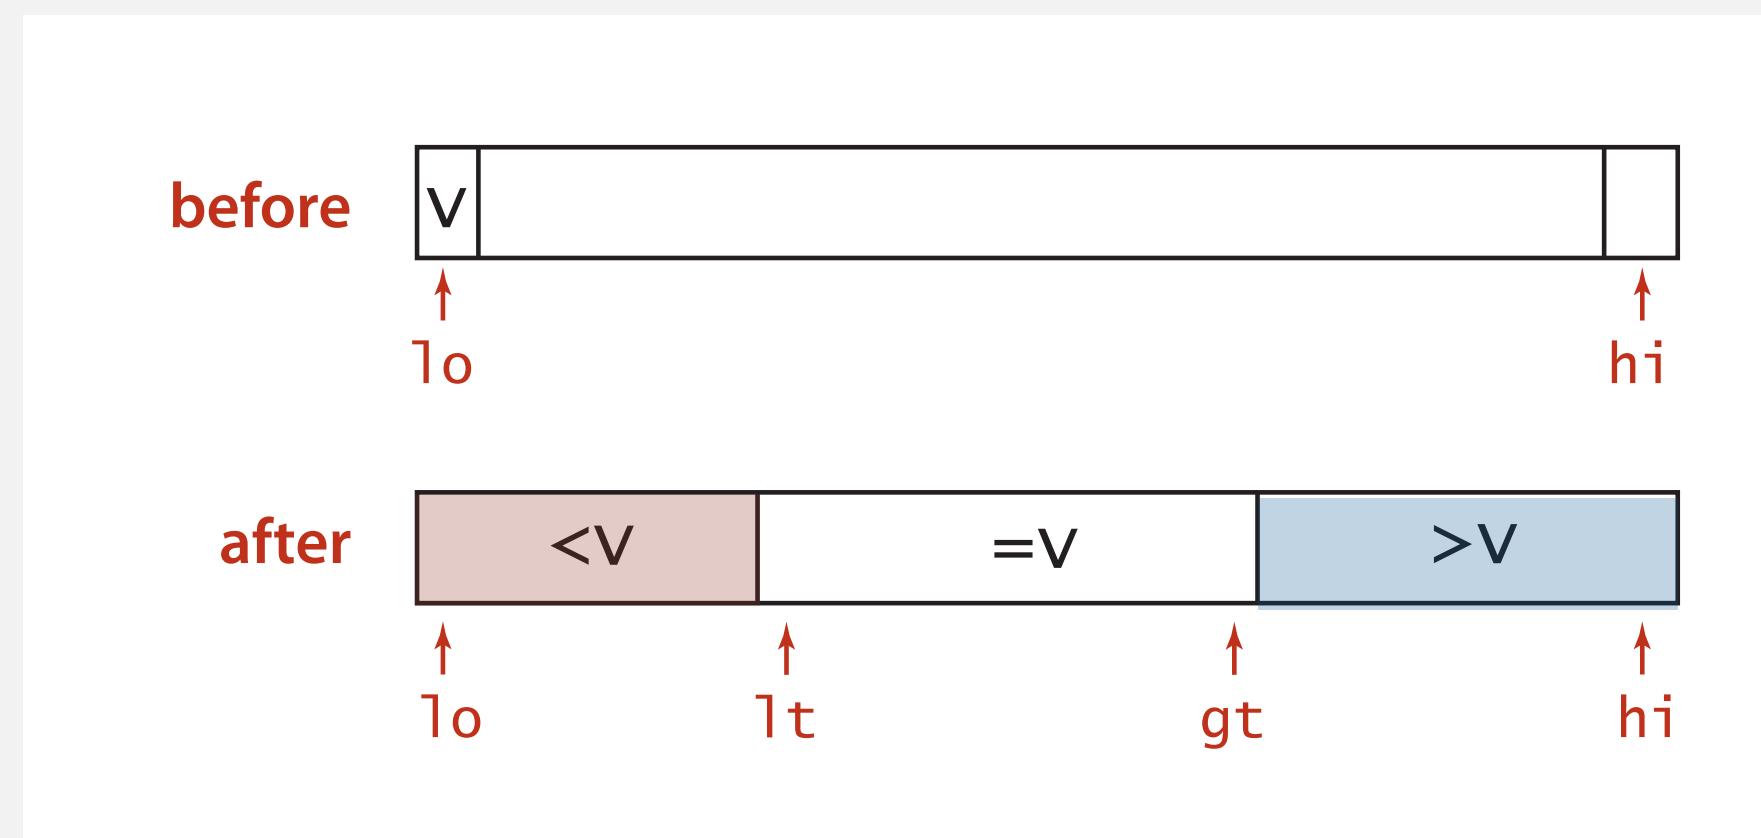
Performance requirements.

- Exactly n calls to $getColor()$.
- At most n calls to $swap()$.
- $\Theta(1)$ extra space.

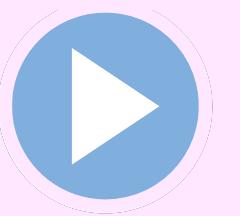
3-way partitioning

Goal. Use pivot $v = a[lo]$ to partition array into **three** parts so that:

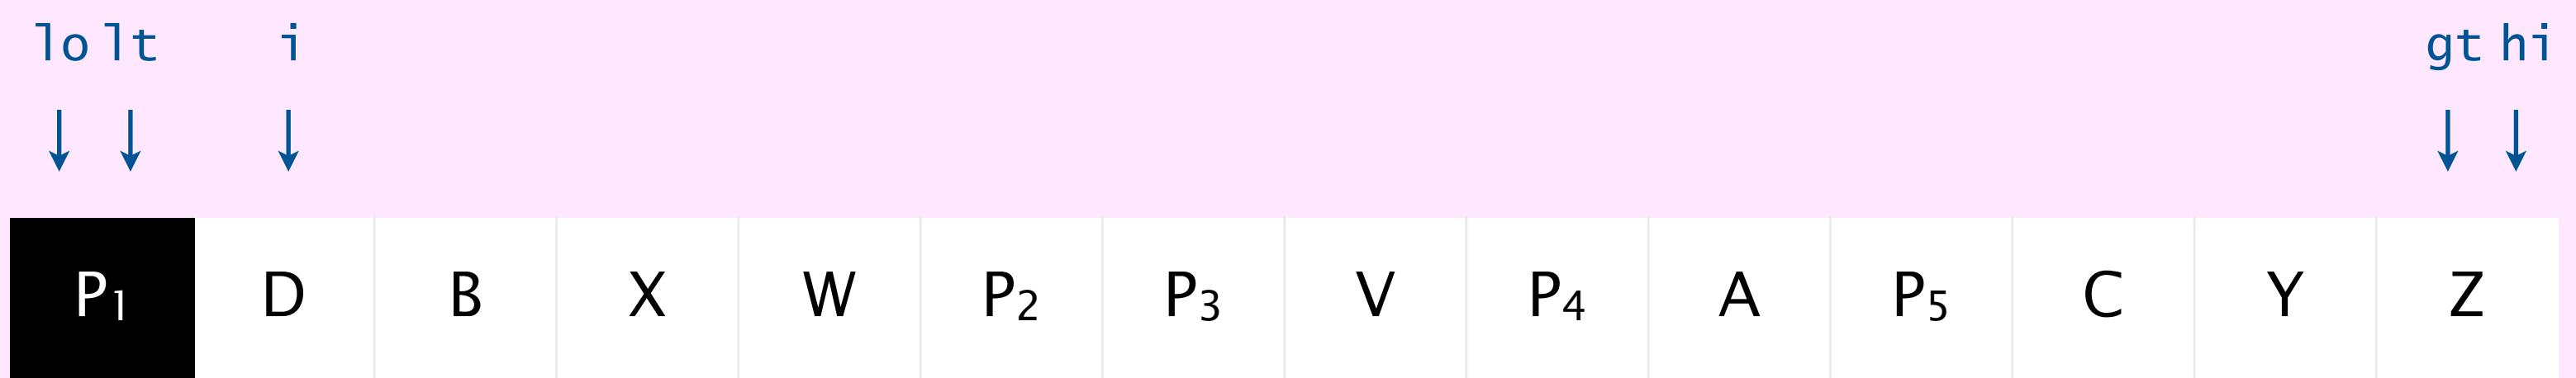
- Red: smaller entries to the left of lt .
- White: equal entries between lt and gt .
- Blue: larger entries to the right of gt .



Dijkstra's 3-way partitioning algorithm: demo



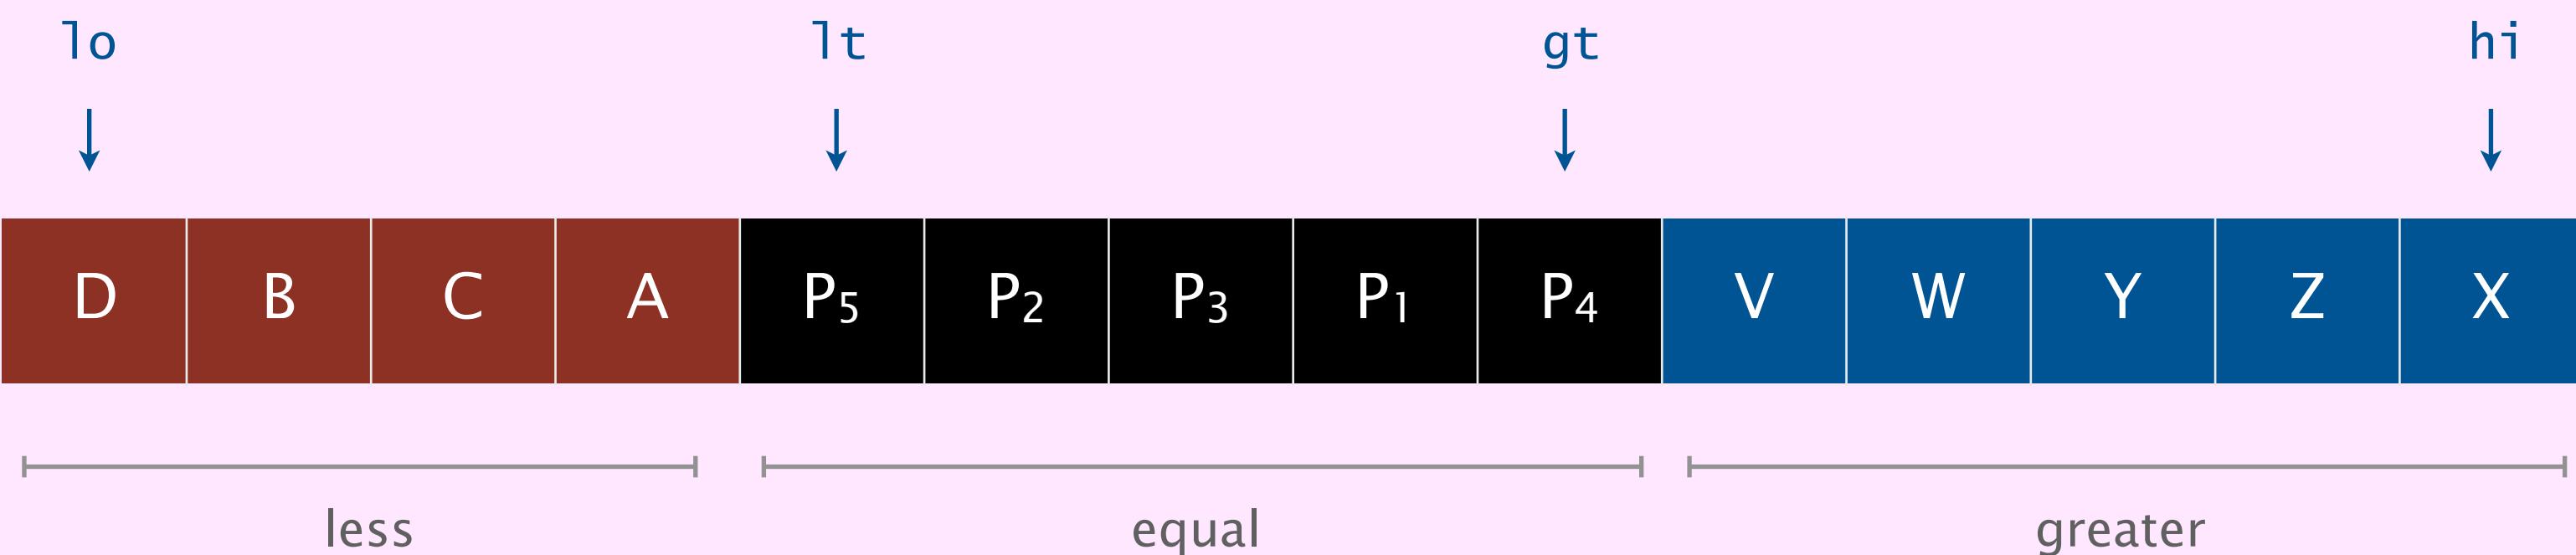
- Let $v = a[lo]$ be pivot.
- Scan i from left to right and compare $a[i]$ to v .
 - less: exchange $a[i]$ with $a[lt]$; increment both lt and i
 - greater: exchange $a[i]$ with $a[gt]$; decrement gt
 - equal: increment i



Dijkstra's 3-way partitioning algorithm: demo

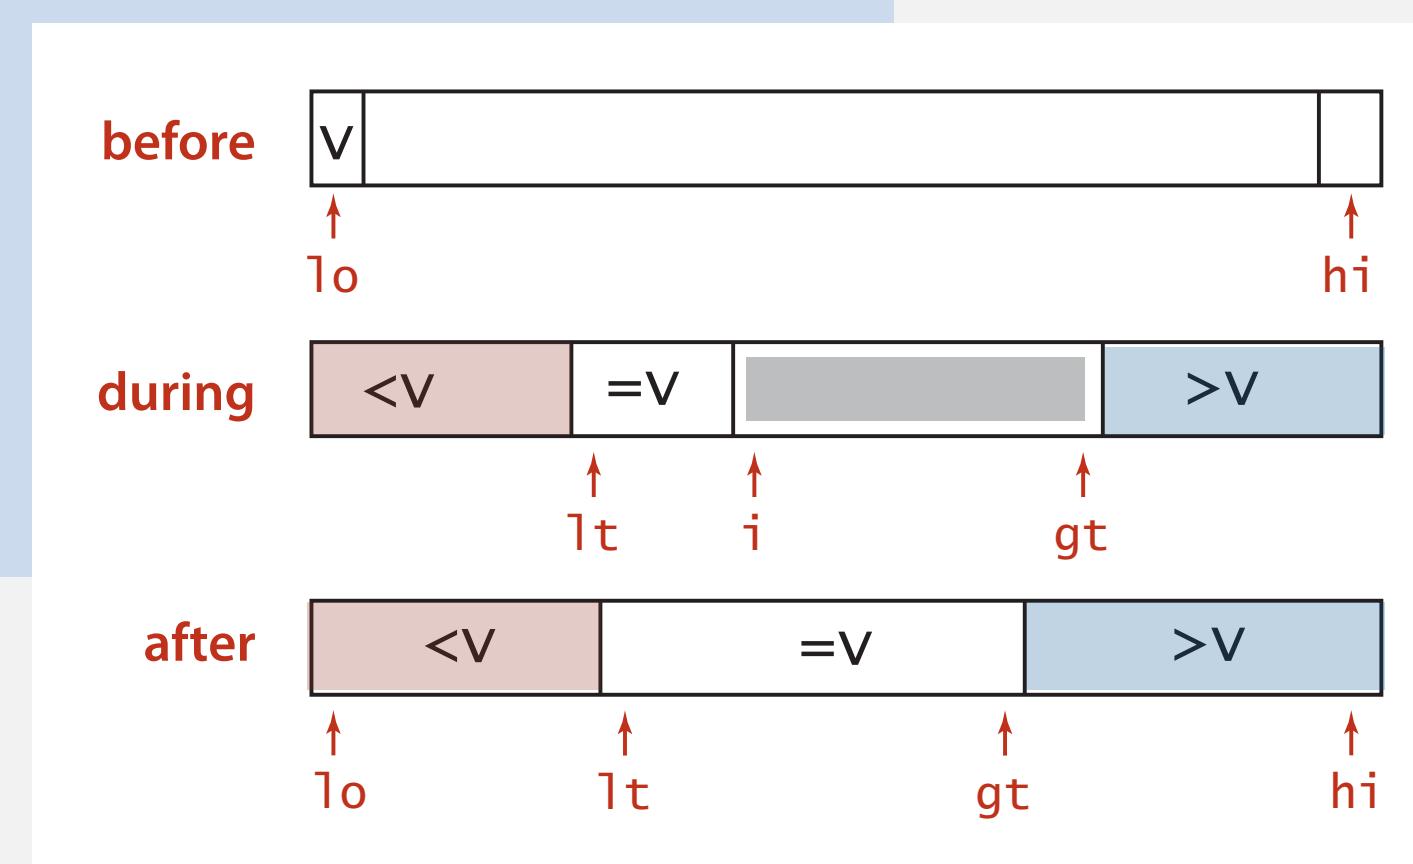


- Let $v = a[lo]$ be pivot.
- Scan i from left to right and compare $a[i]$ to v .
 - less: exchange $a[i]$ with $a[lt]$; increment both lt and i
 - greater: exchange $a[i]$ with $a[gt]$; decrement gt
 - equal: increment i



3-way quicksort: Java implementation

```
private static void sort(Comparable[] a, int lo, int hi)
{
    if (hi <= lo) return;
    int lt = lo, gt = hi;
    Comparable v = a[lo];
    int i = lo + 1;
    while (i <= gt)
    {
        int cmp = a[i].compareTo(v);
        if      (cmp < 0) exch(a, lt++, i++);
        else if (cmp > 0) exch(a, i, gt--);
        else               i++;
    }
    sort(a, lo, lt - 1);
    sort(a, gt + 1, hi);
}
```

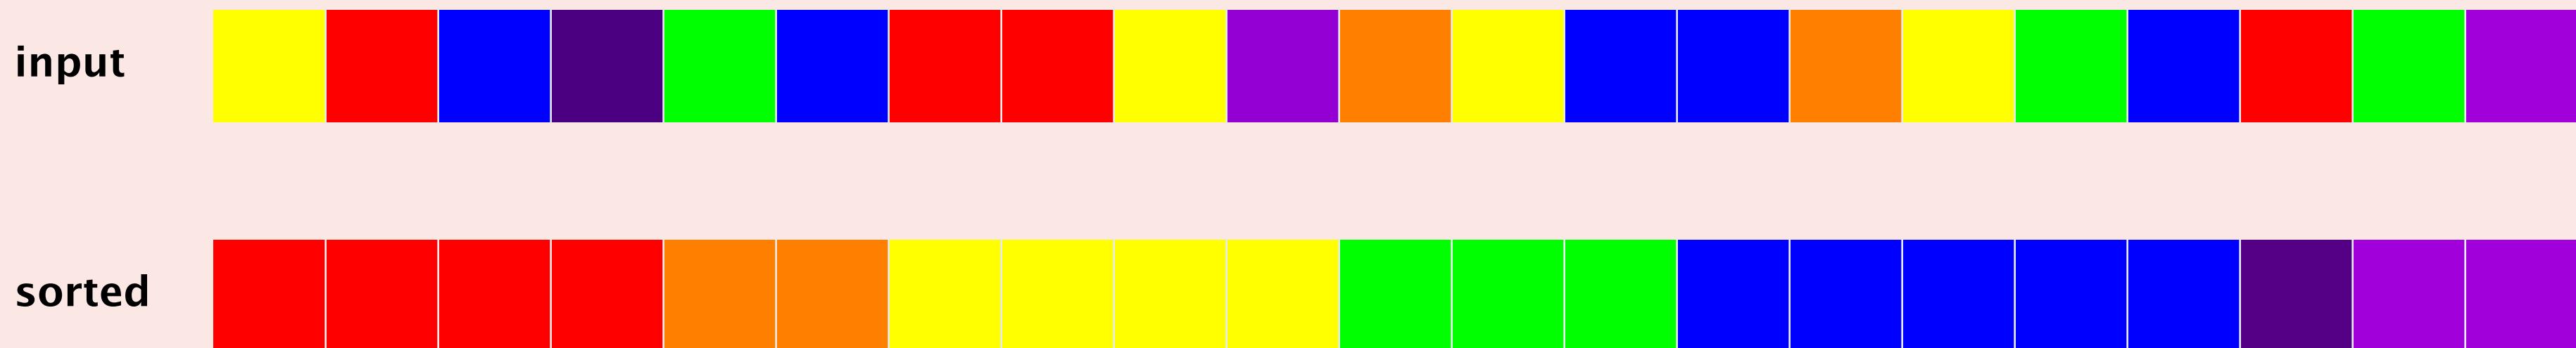




Quicksort quiz 5

What is the worst-case number of compares to 3-way quicksort an array of length n containing only 7 distinct values?

- A. $\Theta(n)$
- B. $\Theta(n \log n)$
- C. $\Theta(n^2)$
- D. $\Theta(n^7)$



Sorting summary

	inplace?	stable?	best	average	worst	remarks
selection	✓		$\frac{1}{2} n^2$	$\frac{1}{2} n^2$	$\frac{1}{2} n^2$	n exchanges
insertion	✓	✓	n	$\frac{1}{4} n^2$	$\frac{1}{2} n^2$	use for small n or partially sorted arrays
merge		✓	$\frac{1}{2} n \log_2 n$	$n \log_2 n$	$n \log_2 n$	$\Theta(n \log n)$ guarantee; stable
timsort		✓	n	$n \log_2 n$	$n \log_2 n$	improves mergesort when pre-existing order
quick	✓		$n \log_2 n$	$2 n \ln n$	$\frac{1}{2} n^2$	$\Theta(n \log n)$ probabilistic guarantee; fastest in practice
3-way quick	✓		n	$2 n \ln n$	$\frac{1}{2} n^2$	improves quicksort when duplicate keys
?	✓	✓	n	$n \log_2 n$	$n \log_2 n$	holy sorting grail

number of compares to sort an array of n elements



2.3 QUICKSORT

- ▶ *quicksort*
- ▶ *selection*
- ▶ *duplicate keys*
- ▶ ***system sorts***

Sorting applications

Sorting algorithms are essential in a broad variety of applications:

- Sort a list of names.
- Organize an MP3 library.
- Display Google PageRank results.
- List RSS feed in reverse chronological order.

obvious applications

- Find the median.
- Identify statistical outliers.
- Binary search in a database.
- Find duplicates in a mailing list.

problems become easy once
items are in sorted order

- Data compression.
- Computer graphics.
- Computational biology.
- Load balancing on a parallel computer.

non-obvious applications

...

Engineering a system sort (in 1993)

Bentley–McIlroy quicksort.

- Cutoff to insertion sort for small subarrays.
- Pivot selection: median of 3 or Tukey's ninther.
- Partitioning scheme: Bentley–McIlroy 3-way partitioning.

sample 9 items

similar to Dijkstra 3-way partitioning

(but fewer exchanges when not many equal keys)

Engineering a Sort Function

JON L. BENTLEY

M. DOUGLAS McILROY

AT&T Bell Laboratories, 600 Mountain Avenue, Murray Hill, NJ 07974, U.S.A.

SUMMARY

We recount the history of a new `qsort` function for a C library. Our function is clearer, faster and more robust than existing sorts. It chooses partitioning elements by a new sampling scheme; it partitions by a novel solution to Dijkstra's Dutch National Flag problem; and it swaps efficiently. Its behavior was assessed with timing and debugging testbeds, and with a program to certify performance. The design techniques apply in domains beyond sorting.

In the wild. C, C++, Java 6,

A Java mailing list post (Yaroslavskiy, September 2009)

Replacement of quicksort in java.util.Arrays with new dual-pivot quicksort

Hello All,

I'd like to share with you new **Dual-Pivot Quicksort** which is faster than the known implementations (theoretically and experimental). I'd like to propose to replace the JDK's Quicksort implementation by new one.

...

The new Dual-Pivot Quicksort uses ***two*** pivots elements in this manner:

1. Pick an elements P1, P2, called pivots from the array.
2. Assume that P1 <= P2, otherwise swap it.
3. Reorder the array into three parts: those less than the smaller pivot, those larger than the larger pivot, and in between are those elements between (or equal to) the two pivots.
4. Recursively sort the sub-arrays.

The invariant of the Dual-Pivot Quicksort is:

[< P1 | P1 <= & <= P2 } > P2]

...

Another Java mailing list post (Yaroslavskiy-Bloch-Bentley)

Replacement of quicksort in java.util.Arrays with new dual-pivot quicksort

Date: Thu, 29 Oct 2009 11:19:39 +0000

Subject: Replace quicksort in java.util.Arrays with dual-pivot implementation

Changeset: b05abb410c52

Author: alanb

Date: 2009-10-29 11:18 +0000

URL: <http://hg.openjdk.java.net/jdk7/tl/jdk/rev/b05abb410c52>

6880672: Replace quicksort in java.util.Arrays with dual-pivot implementation

Reviewed-by: jjb

Contributed-by: vladimir.yaroslavskiy at sun.com, joshua.bloch at google.com,
jbentley at avaya.com

! src/share/classes/java/util/Arrays.java

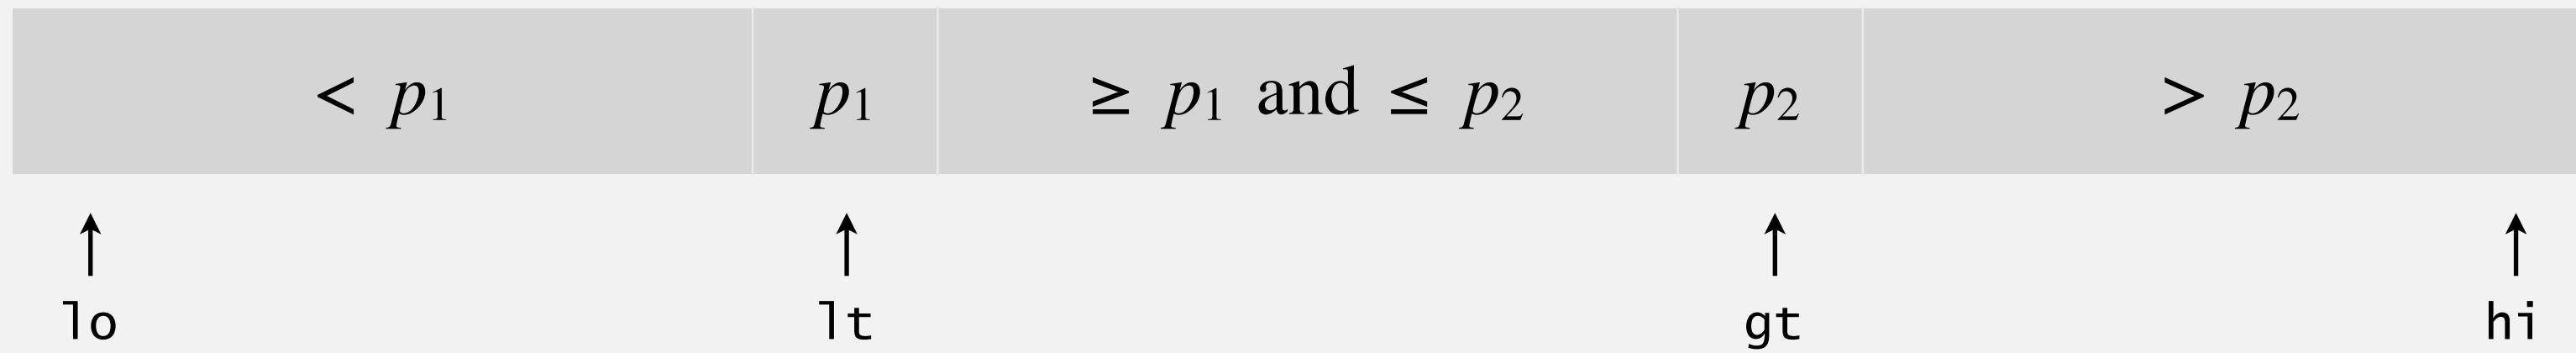
+ src/share/classes/java/util/DualPivotQuicksort.java

<https://mail.openjdk.java.net/pipermail/compiler-dev/2009-October.txt>

Dual-pivot quicksort

Use two pivots p_1 and p_2 and partition into three subarrays:

- Keys less than p_1 .
- Keys between p_1 and p_2 .
- Keys greater than p_2 .



Recursively sort three subarrays (skip middle subarray if $p_1 = p_2$).

degenerates to Dijkstra's 3-way partitioning

In the wild. Java 8, Python unstable sort, Android, ...

SYSTEM SORT



Suppose you are the lead architect of a new programming language.

Which sorting algorithm(s) would you use for the system sort? Defend your answer.

System sorts in Java 8 and Java 11

`Arrays.sort()` and `Arrays.parallelSort()`.

- Has one method for `Comparable` objects.
- Has an overloaded method for each primitive type.
- Has an overloaded method for use with a `Comparator`.
- Has overloaded methods for sorting subarrays.



Algorithms.

- Timsort for reference types.
- Dual-pivot quicksort for primitive types.
- Parallel mergesort for `Arrays.parallelSort()`.

Q. Why use different algorithms for primitive and reference types?

Bottom line. Use the system sort!

© Copyright 2021 Robert Sedgewick and Kevin Wayne