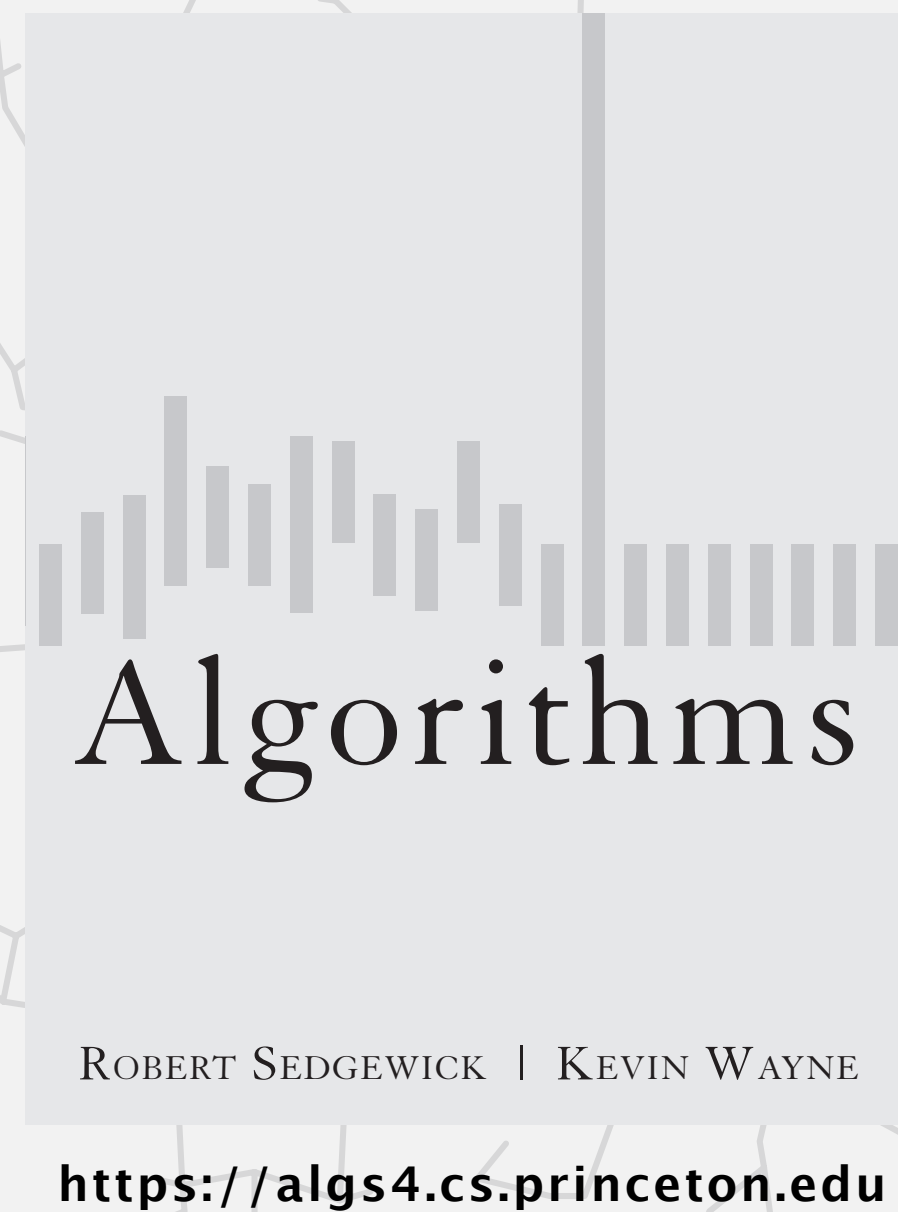




<https://algs4.cs.princeton.edu>

2.4 PRIORITY QUEUES

- ▶ *APIs*
- ▶ *elementary implementations*
- ▶ *binary heaps*
- ▶ *heapsort*
- ▶ *event-driven simulation* ← see videos



2.4 PRIORITY QUEUES


- ▶ *APIs*
- ▶ *elementary implementations*
- ▶ *binary heaps*
- ▶ *heapsort*
- ▶ *event-driven simulation*

Collections

A **collection** is a data type that stores a group of items.

data type	core operations	data structure
stack	PUSH, POP	<i>linked list</i> <i>resizing array</i>
queue	ENQUEUE, DEQUEUE	
priority queue	INSERT, DELETE-MAX	<i>binary heap</i>
symbol table	PUT, GET, DELETE	<i>binary search tree</i> <i>hash table</i>
set	ADD, CONTAINS, DELETE	

“ *Show me your code and conceal your data structures, and I shall continue to be mystified. Show me your data structures, and I won’t usually need your code; it’ll be obvious.* ” — *Fred Brooks*



Priority queue

Collections. Insert and remove items. Which item to remove?


Stack. Remove the item most recently added.

Queue. Remove the item least recently added.

Randomized queue. Remove a random item.

Priority queue. Remove the **largest** (or **smallest**) item.

Generalizes: stack, queue, randomized queue.



triage in an emergency room
(priority = urgency of wound/illness)

operation	argument	return value
insert	P	
insert	Q	
insert	E	
remove max		Q
insert	X	
insert	A	
insert	M	
remove max		X
insert	P	
insert	L	
insert	E	
remove max		P

Max-oriented priority queue API

Requirement. Must insert keys of the same (generic) type; keys must be Comparable.

“bounded type parameter”
↙

```
public class MaxPQ<Key extends Comparable<Key>>
```

MaxPQ()	<i>create an empty priority queue</i>
void insert(Key v)	<i>insert a key</i>
Key delMax()	<i>return and remove a largest key</i>
boolean isEmpty()	<i>is the priority queue empty?</i>
Key max()	<i>return a largest key</i>
int size()	<i>number of entries in the priority queue</i>

Note. Duplicate keys allowed; delMax() removes and returns any maximum key.

Min-oriented priority queue API

Analogous to MaxPQ.

```
public class MinPQ<Key extends Comparable<Key>>
```

<code>MinPQ()</code>	<i>create an empty priority queue</i>
----------------------	---------------------------------------

<code>void insert(Key v)</code>	<i>insert a key</i>
---------------------------------	---------------------

<code>Key delMin()</code>	<i>return and remove a smallest key</i>
---------------------------	---

<code>boolean isEmpty()</code>	<i>is the priority queue empty?</i>
--------------------------------	-------------------------------------

<code>Key min()</code>	<i>return a smallest key</i>
------------------------	------------------------------

<code>int size()</code>	<i>number of entries in the priority queue</i>
-------------------------	--

Warmup client. Sort a stream of integers from standard input.

Priority queue: applications


- Event-driven simulation. [customers in a line, colliding particles]
- Discrete optimization. [bin packing, scheduling]
- Artificial intelligence. [A* search]
- Computer networks. [web cache]
- Data compression. [Huffman codes]
- Operating systems. [load balancing, interrupt handling]
- Graph searching. [Dijkstra's algorithm, Prim's algorithm]
- Number theory. [sum of powers]
- Spam filtering. [Bayesian spam filter]
- Statistics. [online median in data stream]



priority = length of
best known path

8	4	7
1	5	6
3	2	

priority = "distance"
to goal board



priority = event time




2.4 PRIORITY QUEUES

- ▶ *APIs*
- ▶ *elementary implementations*
- ▶ *binary heaps*
- ▶ *heapsort*
- ▶ *event-driven simulation*

Priority queue: elementary implementations

Unordered list. Store keys in a linked list.



Performance. INSERT takes $\Theta(1)$ time; DELETE-MAX takes $\Theta(n)$ time.

Priority queue: elementary implementations

Ordered array. Store keys in an array in ascending (or descending) order.




ordered array implementation of a MaxPQ



What are the worst-case running times for INSERT and DELETE-MAX, respectively, in a MaxPQ implemented with an **ordered array**?

ignore array resizing

- A. $\Theta(1)$ and $\Theta(n)$
- B. $\Theta(1)$ and $\Theta(\log n)$
- C. $\Theta(\log n)$ and $\Theta(1)$
- D. $\Theta(n)$ and $\Theta(1)$



ordered array implementation of a MaxPQ

Priority queue: implementations cost summary

Elementary implementations. Either INSERT or DELETE-MAX takes $\Theta(n)$ time.

implementation	INSERT	DELETE-MAX	MAX
unordered list	1	n	n
ordered array	n	1	1
goal	$\log n$	$\log n$	$\log n$

order of growth of running time for priority queue with n items

Challenge. Implement **both** core operations efficiently.

Solution. “**Somewhat-ordered**” array.




2.4 PRIORITY QUEUES

- ▶ *API*
- ▶ *elementary implementations*
- ▶ *binary heaps*
- ▶ *heapsort*
- ▶ *event-driven simulation*

Complete binary tree

Binary tree. Empty **or** node with links to two disjoint binary trees (left and right subtrees).

Complete tree. Every level (except possibly the last) is completely filled; the last level is filled from left to right.



complete binary tree with $n = 16$ nodes (height = 4)


Property. Height of complete binary tree with n nodes is $\lfloor \log_2 n \rfloor$.

Pf. As you successively add nodes, height increases (by 1) only when n is a power of 2.

A complete binary tree in nature (of height 4)



A complete binary tree (of height 15)



Binary heap: representation


Binary heap. Array representation of a heap-ordered complete binary tree.

Heap-ordered tree.


- Keys in nodes.
- Child's key no larger than parent's key.

Array representation.

- Indices start at 1.
- Take nodes in **level order**.
- No explicit links!



	0	1	2	3	4	5	6	7	8	9	10	11
a[]	–	T	S	R	P	N	O	A	P	I	H	G




Consider the node at index k in a binary heap. Which Java expression produces the index of its parent?

A. $(k - 1) / 2$

B. $k / 2$

C. $(k + 1) / 2$

D. $2 * k$




	0	1	2	3	4	5	6	7	8	9	10	11
a[]	-	T	S	R	P	N	O	A	E	I	H	G

Binary heap: properties


Proposition. Largest key is at index 1, which is root of binary tree.

Proposition. Can use array indices to move up or down tree.

- Parent of key at index k is at index $k/2$.
- Children of key at index k are at indices $2*k$ and $2*k + 1$.





	0	1	2	3	4	5	6	7	8	9	10	11
a[]	–	T	S	R	P	N	O	A	P	I	H	G



Insert. Add node at end, then **swim** it up.

Remove the maximum. Exchange root with node at end, then **sink** it down.

heap ordered



Binary heap: promotion


Scenario. Key in node becomes **larger** than key in parent's node.

To eliminate the violation:

- Exchange key in child node with key in parent node.
- Repeat until heap order restored.

```
private void swim(int k)
{
    while (k > 1 && less(k/2, k))
    {
        exch(k, k/2);
        k = k/2;
    }
}
```

parent of node at k is at k/2



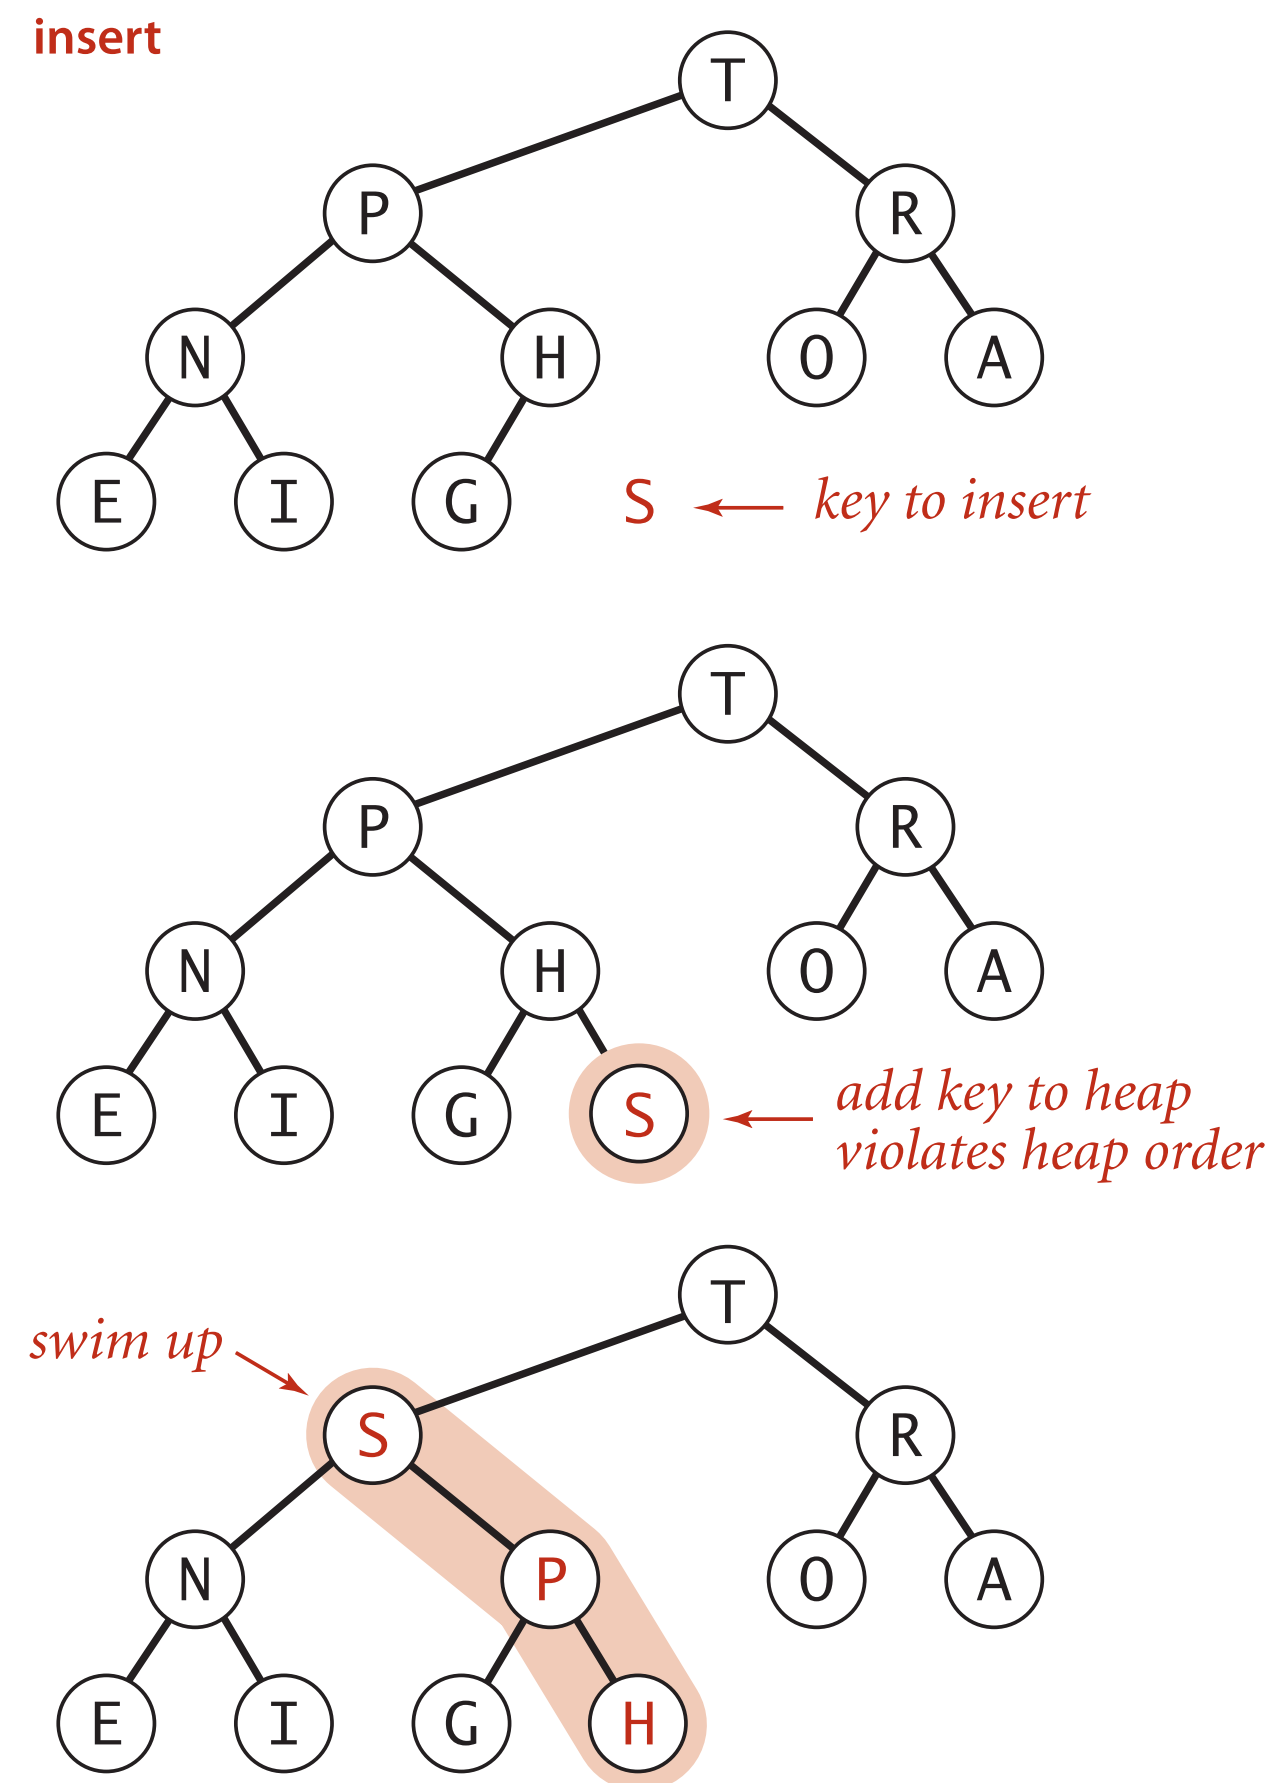
Peter principle. Node promoted to level of incompetence.

Binary heap: insertion

Insert. Add node at end in bottom level; then, swim it up.

Cost. At most $1 + \log_2 n$ compares.

```
public void insert(Key x)
{
    pq[++n] = x;
    swim(n);
}
```



Binary heap: demotion

Scenario. Key in node becomes **smaller** than one (or both) of keys in childrens' nodes.

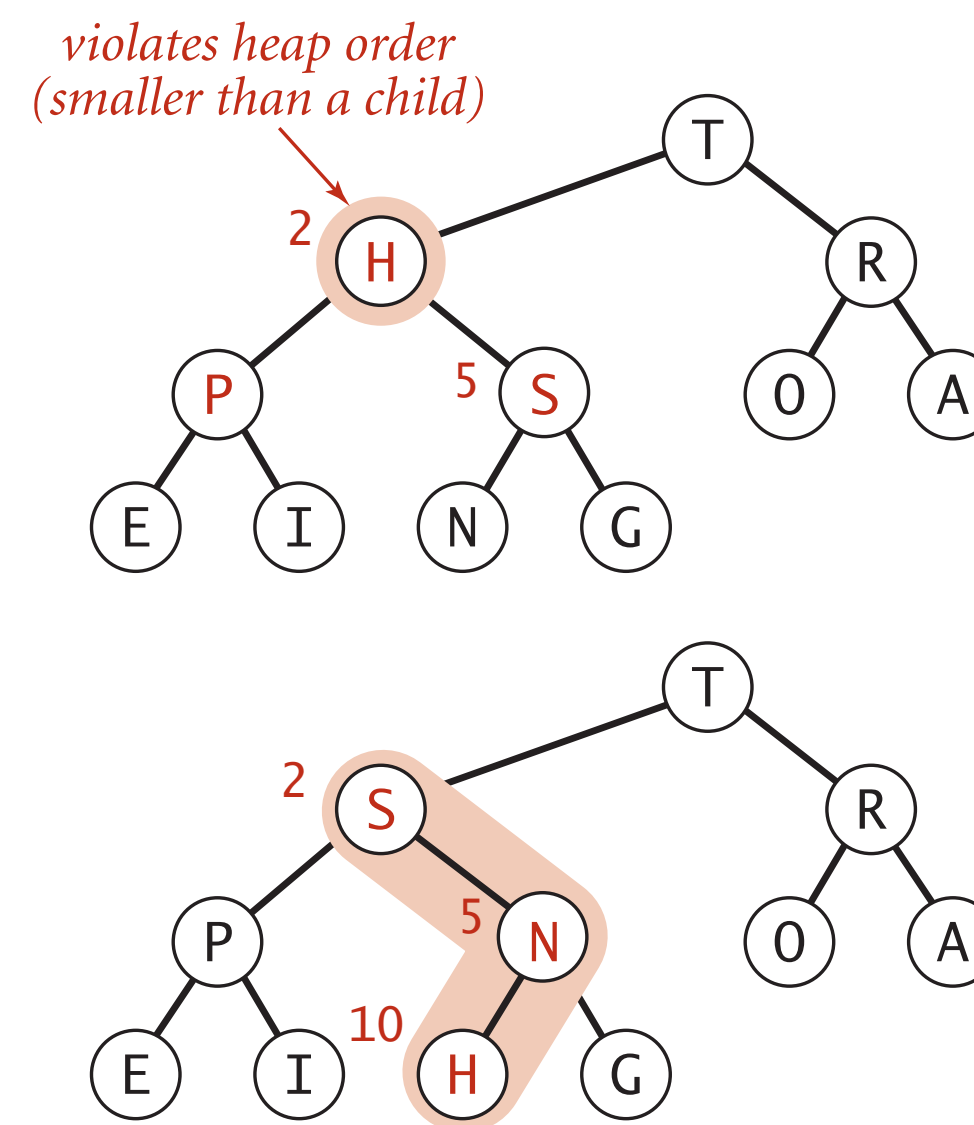
To eliminate the violation:

why not smaller child?

- Exchange key in parent node with key in larger child's node.
- Repeat until heap order restored.

```
private void sink(int k)
{
    while (2*k <= n)
    {
        int j = 2*k;
        if (j < n && less(j, j+1)) j++;
        if (!less(k, j)) break;
        exch(k, j);
        k = j;
    }
}
```

children of node at k
are at 2*k and 2*k+1



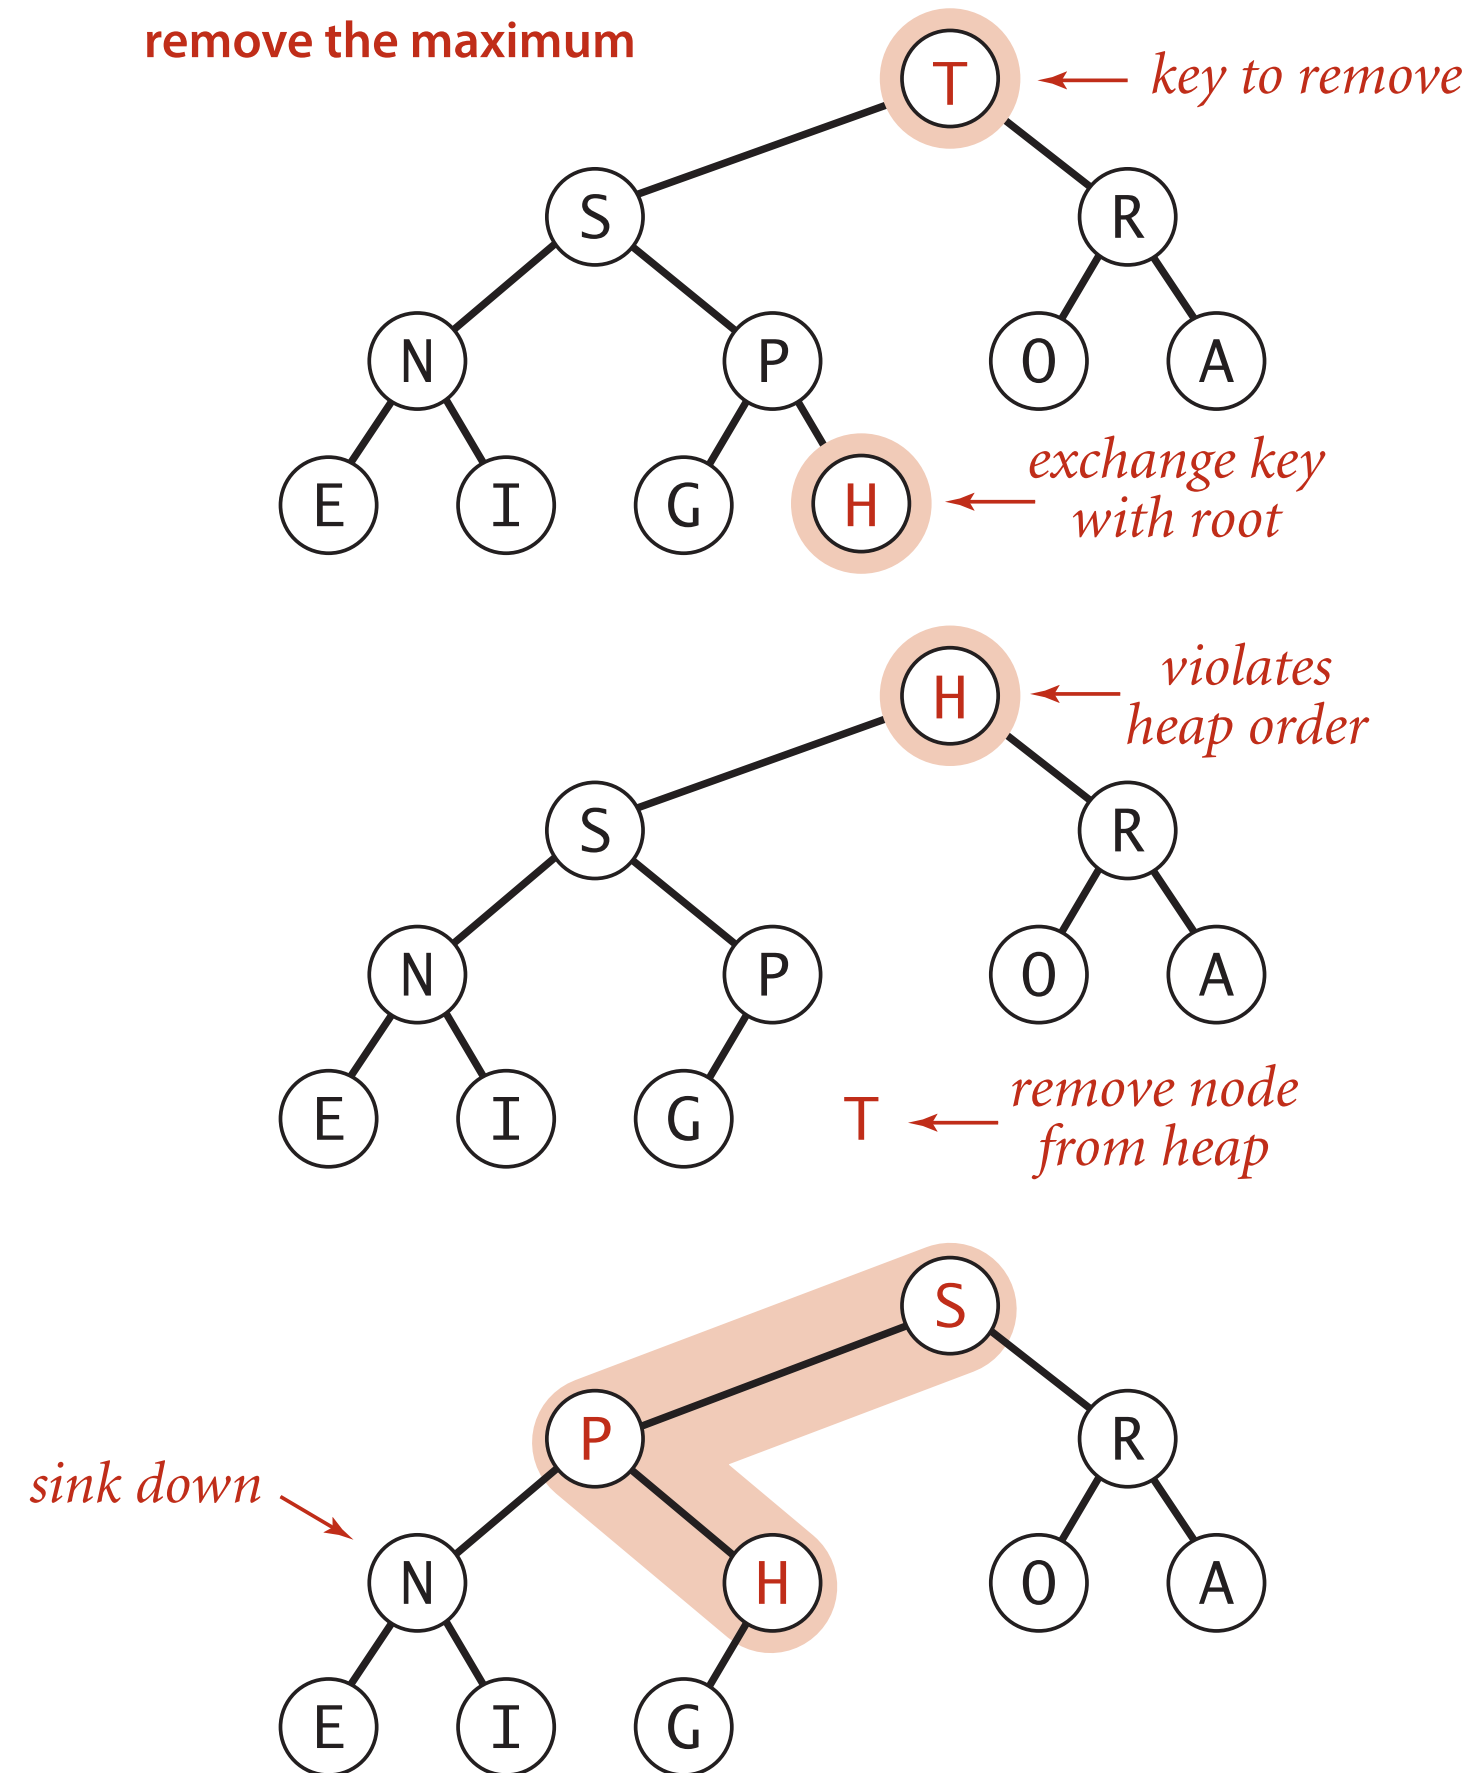
Power struggle. Better subordinate promoted.

Binary heap: delete the maximum

Delete max. Exchange root with node at end; then, sink it down.

Cost. At most $2 \log_2 n$ compares.

```
public Key delMax()
{
    Key max = pq[1];
    exch(1, n--);
    sink(1);
    pq[n+1] = null; ← prevent loitering
    return max;
}
```



Binary heap: Java implementation

```
public class MaxPQ<Key extends Comparable<Key>>
{
```

```
    private Key[] a;
    private int n;
```

```
    public MaxPQ(int capacity)
    { a = (Key[]) new Comparable[capacity+1]; }
```

← fixed capacity
(for simplicity)

```
    public boolean isEmpty()
    { return n == 0; }
    public void insert(Key key) // see previous code
    public Key delMax()         // see previous code
```

← PQ ops

```
    private void swim(int k) // see previous code
    private void sink(int k) // see previous code
```

← heap helper functions

```
    private boolean less(int i, int j)
    { return a[i].compareTo(a[j]) < 0; }
    private void exch(int i, int j)
    { Key temp = a[i]; a[i] = a[j]; a[j] = temp; }
```

← array helper functions

```
}
```

Priority queue: implementations cost summary

Goal. Implement both INSERT and DELETE-MAX in $\Theta(\log n)$ time.

implementation	INSERT	DELETE-MAX	MAX
unordered list	1	n	n
ordered array	n	1	1
goal	$\log n$	$\log n$	1

order of growth of running time for priority queue with n items

Binary heap: considerations

Underflow and overflow.

- Underflow: throw exception if deleting from empty PQ.
- Overflow: add no-arg constructor and use resizing array.

leads to $O(\log n)$
amortized time per op
(how to make worst case?)

Minimum-oriented priority queue.

- Replace `less()` with `greater()`.
- Implement `greater()`.

Other operations.

- Remove an arbitrary item.
- Change the priority of an item.


← can implement efficiently with `sink()` and `swim()`
[stay tuned for Prim/Dijkstra]

Immutability of keys.

- Assumption: client does not change keys while they're on the PQ.
- Best practice: use immutable keys.

← immutable in Java: `String`, `Integer`, `Double`, ...

PRIORITY QUEUE WITH DELETE-RANDOM



Goal. Design an efficient data structure to support the following API:

- INSERT: insert a key.
- DELETE-MAX: return and remove a largest key.
- **SAMPLE:** return a random key.
- **DELETE-RANDOM:** return and remove a random key.




Multiway heaps

Multiway heaps.


- Complete d -way tree.
- Child's key no larger than parent's key.

Property. Height of complete d -way tree on n nodes is $\sim \log_d n$.

Property. Children of key at index k are at indices $3k - 1$, $3k$, and $3k + 1$.




3-way heap



In the worst case, how many compares to **INSERT** and **DELETE-MAX** in a d -way heap as function of both n and d ?


- A. $\sim \log_d n$ and $\sim \log_d n$
- B. $\sim \log_d n$ and $\sim d \log_d n$
- C. $\sim d \log_d n$ and $\sim \log_d n$
- D. $\sim d \log_d n$ and $\sim d \log_d n$



Priority queue: implementation cost summary


implementation	INSERT	DELETE-MAX	MAX	
unordered list	1	n	n	
ordered array	n	1	1	
binary heap	$\log n$	$\log n$	1	
d-ary heap	$\log_d n$	$d \log_d n$	1	← sweet spot: $d = 4$
Fibonacci	1	$\log n$	1	← see COS 423
impossible	1	1	1	← why impossible?

order-of-growth of running time for priority queue with n items



2.4 PRIORITY QUEUES

- ▶ *APIs*
- ▶ *elementary implementations*
- ▶ *binary heaps*
- ▶ ***heapsort***
- ▶ *event-driven simulation*



What are the properties of this sorting algorithm?

```
public void sort(String[] a)
{
    int n = a.length;
    MinPQ<String> pq = new MinPQ<String>();

    for (int i = 0; i < n; i++)
        pq.insert(a[i]);

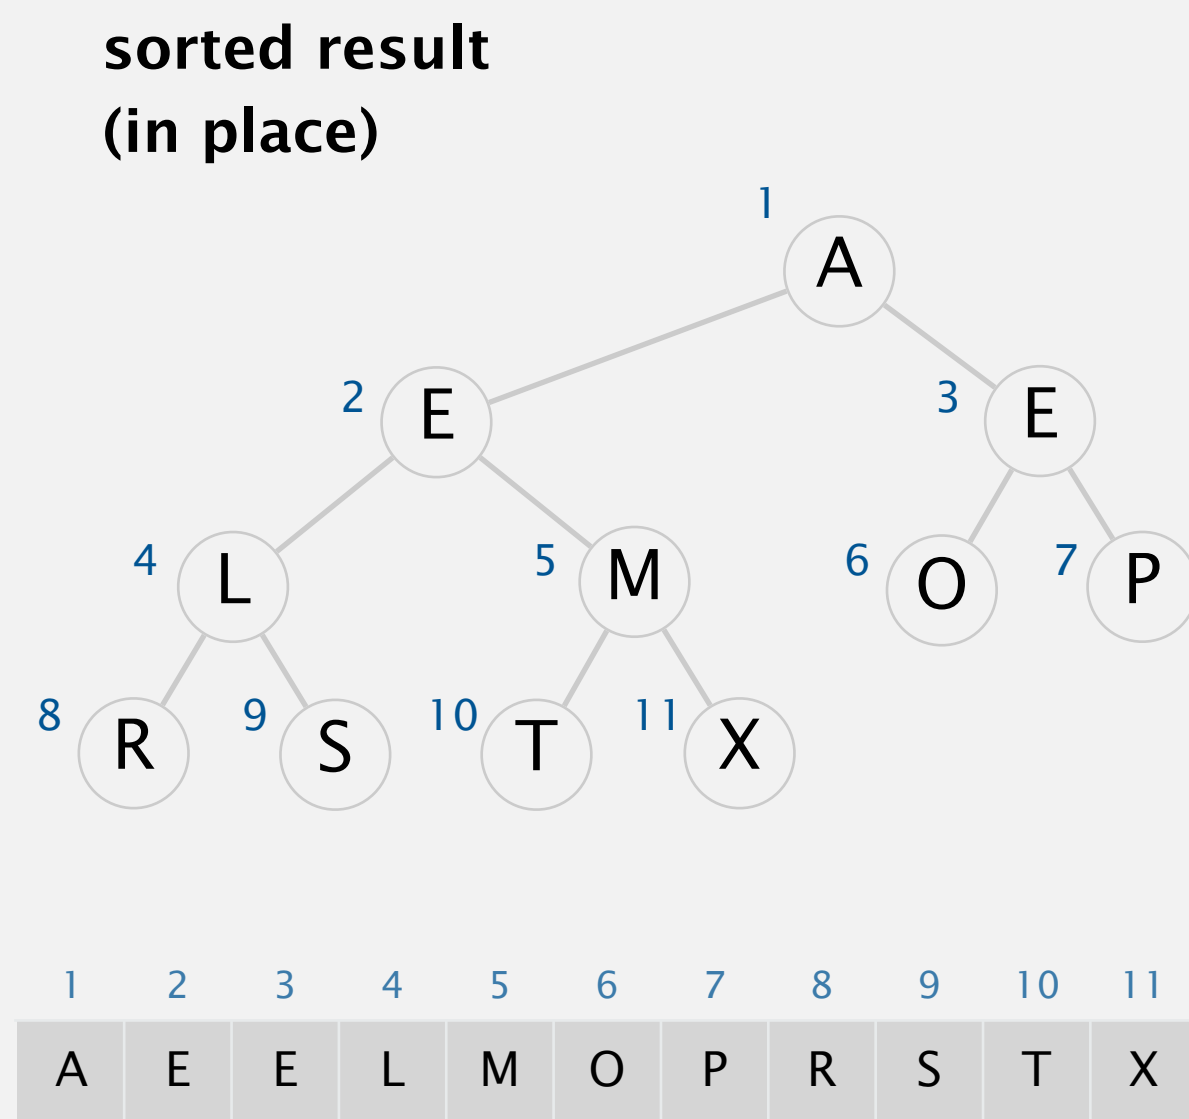


    for (int i = 0; i < n; i++)
        a[i] = pq.delMin();
}
```

- A. $\Theta(n \log n)$ compares in the worst case.
- B. In-place.
- C. Stable.
- D. *All of the above.*

Heapsort

Basic plan for in-place sort.

- View input array as a complete binary tree. ← we'll assume 1-indexed for now
- Heap construction: build a **max-oriented** heap with all n keys.
- Sortdown: repeatedly remove the maximum key.



Heapsort: top-down heap construction


Top-down heap construction. Insert keys into a max heap, one at a time.

```
for (int k = 1; k <= n; k++)  
    swim(a, k);
```

Invariants. After calling `swim(a, k)`,


- `a[1..k]` is a max heap.
- `a[k+1..n]` are untouched.

before call to `swim(a, k)`
`k = 6`



1	2	3	4	5	6	7	8	9	10	11
T	S	R	O	E	X	A	M	P	L	E

after call to `swim(a, k)`
`k = 6`



1	2	3	4	5	6	7	8	9	10	11
X	S	T	O	E	R	A	M	P	L	E

Heapsort: sortdown

Second pass.

- Remove the maximum, one at a time.
- Leave in array (instead of nulling out).


Invariants. After calling `sink(a, 1, k)`,

- `a[1..k-1]` is a max heap.
- `a[k..n]` are in final sorted order.

```
int k = n;  
while (k > 1)  
{  
    exch(a, 1, k--);  
    sink(a, 1, k);  
}
```


← delete-max
(but leave in array)

before call to `sink(a, 1, k)`
`k = 7`



1	2	3	4	5	6	7	8	9	10	11
P	O	E	M	L	E	A	R	S	T	X

after call to `sink(a, 1, k)`
`k = 7`



1	2	3	4	5	6	7	8	9	10	11
O	M	E	A	L	E	P	R	S	T	X

Heapsort: Java implementation

```
public class HeapTopDown
{
    public static void sort(Comparable[] a)
    {
        // top-down heap construction
        int n = a.length;
        for (int k = 1; k <= n; k--)
            swim(a, k);

        // sortdown
        int k = n;
        while (k > 1)
        {
            exch(a, 1, k--);
            sink(a, 1, k);
        }
    }

    ...
}
```

<https://algs4.cs.princeton.edu/24pq/HeapTopDown.java.html>

```
private static void sink(Comparable[] a, int k, int n)
{ /* as before */ }

private static void swim(Comparable[] a, int k)
{ /* as before */ }

private static boolean less(Comparable[] a, int i, int j)
{ /* as before */ }

private static void exch(Object[] a, int i, int j)
{ /* as before */ }
```

but make static
(and pass arguments)

but convert from 1-based
indexing to 0-base indexing

Heapsort: mathematical analysis


Proposition. Heapsort uses only $\Theta(1)$ extra space.

Proposition. Heapsort makes $\leq 3 n \log_2 n$ compares (and $\leq 2 n \log_2 n$ exchanges).

- Top-down heap construction: $\log_2 1 + \log_2 2 + \dots + \log_2 n = \log_2(n!) \sim n \log_2 n$ compares.
- Sortdown: $2 (\log_2 1 + \log_2 2 + \dots + \log_2 n) \sim 2n \log_2 n$ compares.

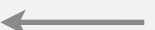
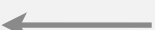
Bottom-up heap construction. [see book] Successively building larger heap from smaller ones.

Proposition. Makes $\leq 2 n$ compares (and $\leq n$ exchanges).




Heapsort: context

Significance. In-place sorting algorithm with $\Theta(n \log n)$ worst-case.

- Mergesort: no, $\Theta(n)$ extra space.  in-place merge possible, not practical
- Quicksort: no, $\Theta(n^2)$ time in worst case.  $\Theta(n \log n)$ worst-case quicksort possible, not practical
- Heapsort: yes!

Bottom line. Heapsort is optimal for both time and space, **but**:


- Inner loop longer than quicksort's.
 - Makes poor use of cache.
 - Not stable.
-  can be improved using advanced caching tricks

Introsort

Goal. As fast as quicksort in practice; $\Theta(n \log n)$ worst case; in place.

Introsort.

- Run quicksort.
- Cutoff to heapsort if function-call stack depth exceeds $2 \log_2 n$.
- Cutoff to insertion sort for $n \leq 16$.



In the wild. C++ STL, Microsoft .NET Framework, Go.

Sorting algorithms: summary

	inplace?	stable?	best	average	worst	remarks
selection	✓		$\frac{1}{2} n^2$	$\frac{1}{2} n^2$	$\frac{1}{2} n^2$	n exchanges
insertion	✓	✓	n	$\frac{1}{4} n^2$	$\frac{1}{2} n^2$	use for small n or partially ordered
merge		✓	$\frac{1}{2} n \log_2 n$	$n \log_2 n$	$n \log_2 n$	$\Theta(n \log n)$ guarantee; stable
timsort		✓	n	$n \log_2 n$	$n \log_2 n$	improves mergesort when pre-existing order
quick	✓		$n \log_2 n$	$2 n \ln n$	$\frac{1}{2} n^2$	$\Theta(n \log n)$ probabilistic guarantee; fastest in practice
3-way quick	✓		n	$2 n \ln n$	$\frac{1}{2} n^2$	improves quicksort when duplicate keys
heap	✓		$3 n$	$2 n \log_2 n$	$2 n \log_2 n$	$\Theta(n \log n)$ guarantee; in-place
?	✓	✓	n	$n \log_2 n$	$n \log_2 n$	holy sorting grail

number of compares to sort an array of n elements

© Copyright 2021 Robert Sedgewick and Kevin Wayne