



<https://algs4.cs.princeton.edu>

## 3.4 HASH TABLES

---

- ▶ *hash functions*
- ▶ *separate chaining*
- ▶ *linear probing*
- ▶ *context*

# Symbol table implementations: summary

---

implementation	guarantee			average case			ordered ops?	key interface
	search	insert	delete	search	insert	delete		
sequential search (unordered list)	$n$	$n$	$n$	$n$	$n$	$n$		equals()
binary search (ordered array)	$\log n$	$n$	$n$	$\log n$	$n$	$n$	✓	compareTo()
BST	$n$	$n$	$n$	$\log n$	$\log n$	$\sqrt{n}$	✓	compareTo()
red-black BST	$\log n$	$\log n$	$\log n$	$\log n$	$\log n$	$\log n$	✓	compareTo()
hashing	$n$	$n$	$n$	1 †	1 †	1 †		equals() hashCode()

† under suitable technical assumptions

Q. Can we do better?

A. Yes, but only with different access to the data.

# Hashing: basic plan

---


Save key-value pairs in a **key-indexed table** (index is a function of the key).


**Hash function:** Mathematical function that maps (hashes) a key to an array index.

**Collision:** Two distinct keys that hash to same index.

**Issue.** Collisions are unavoidable.

- How to limit collisions?
- How to accommodate collisions?





Algorithms

ROBERT SEDGEWICK | KEVIN WAYNE

<https://algs4.cs.princeton.edu>

## 3.4 HASH TABLES

---

- ▶ *hash functions*
- ▶ *separate chaining*
- ▶ *linear probing*
- ▶ *context*


# Designing a hash function

Required properties. [for correctness]


- Deterministic.
- Each key hashes to a table index between  $0$  and  $m - 1$ .

Desirable properties. [for performance]


- Very fast to compute.
- For any subset of  $n$  input keys, each table index gets approximately  $n / m$  keys.



leads to good hash-table performance  
( $m = 10, n = 20$ )



leads to poor hash-table performance  
( $m = 10, n = 20$ )



# Designing a hash function


---

Required properties. [for correctness]

- Deterministic.
- Each key hashes to a table index between  $0$  and  $m - 1$ .



Desirable properties. [for performance]

- Very fast to compute.
- For any subset of  $n$  input keys, each table index gets approximately  $n / m$  keys.



Ex 1. [ $m = 10,000$ ] Last 4 digits of U.S. Social Security number.

Ex 2. [ $m = 10,000$ ] Last 4 digits of phone number.





## Hash tables: quiz 1

Which is the last digit of your **day** of birth?

- A. 0 or 1
- B. 2 or 3
- C. 4 or 5
- D. 6 or 7
- E. 8 or 9





Which is the last digit of your **year of birth?**

- A. 0 or 1
- B. 2 or 3
- C. 4 or 5
- D. 6 or 7
- E. 8 or 9




## Java's hashCode() conventions

---

All Java classes inherit a method `hashCode()`, which returns a 32-bit `int`.

**Requirement.** If `x.equals(y)`, then `(x.hashCode() == y.hashCode())`.

**Highly desirable.** If `!x.equals(y)`, then `(x.hashCode() != y.hashCode())`.



**Customized implementations.** `Integer`, `Double`, `String`, `java.net.URL`, ...

**Legal (but highly undesirable) implementation.** Always return 17.

**User-defined types.** Users are on their own.

# Implementing hashCode(): integers and doubles

## Java library implementations

```
public final class Integer
{
    private final int value;
    ...
    public int hashCode()
    {   return value;   }
}
```

```
public final class Double
{
    private final double value;
    ...
    public int hashCode()
    {
        long bits = doubleToLongBits(value);
        return (int) (bits ^ (bits >>> 32));
    }
}
```

convert to IEEE 64-bit representation;  
xor most significant 32-bits  
with least significant 32-bits

## Implementing hashCode(): arrays

$31x + y$  rule.

- Initialize hash to 1.
- Repeatedly multiply hash by 31 and add next integer in array.

```
public class Arrays
{
    ...

    public static int hashCode(int[] a) {
        if (a == null)
            return 0; ← special case for null

        int hash = 1;
        for (int i = 0; i < a.length; i++)
            hash = 31*hash + a[i];
        return hash;
    }
}
```

**Java library implementation**

prime

$31x + y$  rule

## Implementing hashCode(): user-defined types

```
public final class Transaction
{
    private final String who;
    private final Date when;
    private final double amount;

    public Transaction(String who, Date when, double amount)
    { /* as before */ }

    public boolean equals(Object y)
    { /* as before */ }

    ...
}
```

```
public int hashCode()
{
    int hash = 1;
    hash = 31*hash + who.hashCode();
    hash = 31*hash + when.hashCode();
    hash = 31*hash + ((Double) amount).hashCode();
    return hash;
}
```

for reference types,  
use hashCode()  
  
for primitive types,  
use hashCode()  
of wrapper type

## Implementing hashCode(): user-defined types

---

```
public final class Transaction
{
    private final String who;
    private final Date when;
    private final double amount;

    public Transaction(String who, Date when, double amount)
    { /* as before */ }

    public boolean equals(Object y)
    { /* as before */ }

    ...

    public int hashCode()
    {
        return Objects.hash(who, when, amount); ← shorthand
    }
}
```

## Implementing hashCode()

---

“Standard” recipe for user-defined types.

- Combine each significant field using the  $31x + y$  rule.
- Shortcut 1: use `Objects.hash()` for all fields (except arrays).
- Shortcut 2: use `Arrays.hashCode()` for arrays of primitives.
- Shortcut 3: use `Arrays.deepHashCode()` for arrays of objects.



**Principle.** Every significant field contributes to hash.

**In practice.** Recipe above works reasonably well; used in Java libraries.



## Hash tables: quiz 3

Which Java function maps hashable keys to integers between 0 and  $m-1$  ?

A.


```
private int hash(Key key)
{   return key.hashCode() % m; }
```

B.

```
private int hash(Key key)
{   return Math.abs(key.hashCode()) % m; }
```

C. Both A and B.

D. Neither A nor B.



# Modular hashing

Hash code. An int between  $-2^{31}$  and  $2^{31} - 1$ .

Hash function. An int between 0 and  $m - 1$  (for use as array index).

$m$  typically a prime or a power of 2

```
private int hash(Key key)
{   return key.hashCode() % m; }
```


bug

the remainder operator can evaluate to a negative integer

```
private int hash(Key key)
{   return Math.abs(key.hashCode()) % m; }
```

1-in-a-billion bug

hashCode() of "polygenelubricants" and new Double(-0.0) is  $-2^{31}$



# Modular hashing

Hash code. An int between  $-2^{31}$  and  $2^{31} - 1$ .


Hash function. An int between 0 and  $m - 1$  (for use as array index).

$m$  typically a prime or a power of 2

```
private int hash(Key key)
{   return (key.hashCode() & 0x7fffffff) % m; }
```

correct

discard sign bit



```
private int hash(Key key)
{
    int h = key.hashCode();
    h ^= (h >>> 20) ^ (h >>> 12) ^ (h >>> 7) ^ (h >>> 4);
    return h & (m-1);
}
```

assumes  $m$  is a power of 2


Java 7 (protects against poor quality hashCode())

# Uniform hashing assumption

Uniform hashing assumption. Any key is equally likely to hash to one of  $m$  possible indices.

and independently of other keys

Bins and balls. Toss  $n$  balls uniformly at random into  $m$  bins.



Bad news. [birthday problem]

- In a random group of 23 people, more likely than not that two people share the same birthday.
- Expect two balls in the same bin after  $\sim \sqrt{\pi m / 2}$  tosses.


23.9 when  $m = 365$

# Uniform hashing assumption

Uniform hashing assumption. Any key is equally likely to hash to one of  $m$  possible indices.

and independently of other keys


Bins and balls. Toss  $n$  balls uniformly at random into  $m$  bins.




Good news. [load balancing]

- When  $n \gg m$ , expect most bins to have approximately  $n / m$  balls.
- When  $n = m$ , expect most loaded bin has  $\sim \ln n / \ln \ln n$  balls.

$\text{Binomial}(n, 1 / m)$





Algorithms

ROBERT SEDGEWICK | KEVIN WAYNE

<https://algs4.cs.princeton.edu>

## 3.4 HASH TABLES


---

- ▶ *hash functions*
- ▶ ***separate chaining***
- ▶ *linear probing*
- ▶ *context*

# Collisions

**Collision.** Two distinct keys that hash to the same index.


- Birthday problem  $\Rightarrow$  can't avoid collisions.  $\leftarrow$  unless you have a ridiculous (quadratic) amount of memory
  - Load balancing  $\Rightarrow$  no index gets too many collisions.  
 $\Rightarrow$  ok to scan through all colliding keys.



# Separate-chaining hash table

Use an array of  $m$  linked lists.


- Hash: map key to table index  $i$  between 0 and  $m - 1$ .
- Insert: add key-value pair at front of chain  $i$  (if not already in chain).



# Separate-chaining hash table

Use an array of  $m$  linked lists.

- Hash: map key to table index  $i$  between 0 and  $m - 1$ .
- Insert: add key-value pair at front of chain  $i$  (if not already in chain).
- Search: perform sequential search in chain  $i$ .



# Separate-chaining hash table: Java implementation

```
public class SeparateChainingHashST<Key, Value>
{
    private int m = 128;           // number of chains
    private Node[] st = new Node[m]; // array of chains

    private static class Node
    {
        private Object key;
        private Object val;
        private Node next;
        ...
    }

    private int hash(Key key)
    { /* as before */ }

    public Value get(Key key) {
        int i = hash(key);
        for (Node x = st[i]; x != null; x = x.next)
            if (key.equals(x.key)) return (Value) x.val;
        return null;
    }
}
```

array resizing  
code omitted

no generic array creation  
(declare key and value of type Object)

## Separate-chaining hash table: Java implementation

---

```
public class SeparateChainingHashST<Key, Value>
{
    private int m = 128;                      // number of chains
    private Node[] st = new Node[m]; // array of chains


    private static class Node
    {
        private Object key;
        private Object val;
        private Node next;
        ...
    }

    private int hash(Key key)
    { /* as before */ }

    public void put(Key key, Value val)
    {
        int i = hash(key);
        for (Node x = st[i]; x != null; x = x.next)
            if (key.equals(x.key)) { x.val = val; return; }
        st[i] = new Node(key, val, st[i]);
    }
}
```

## Analysis of separate chaining

Recall load balancing. Under uniform hashing assumption, length of each chain is tightly concentrated around mean =  $n / m$ .



calls to either  
equals() or hashCode()



Consequence. Expected number of probes for search/insert is  $\Theta(n / m)$ .

- $m$  too small  $\Rightarrow$  chains too long.
- $m$  too large  $\Rightarrow$  too many empty chains.
- Typical choice:  $m \sim \frac{1}{4}n \Rightarrow \Theta(1)$  time for search/insert.

$m$  times faster than  
sequential search




# Resizing in a separate-chaining hash table


**Goal.** Average length of chain  $n / m$  is  $\Theta(1)$ .

- Double length  $m$  of array when  $n / m \geq 8$ .
- Halve length  $m$  of array when  $n / m \leq 2$ .
- Note: need to rehash all keys when resizing. ← x.hashCode() does not change;  
but hash(x) typically does

**before resizing ( $n/m = 8$ )**




**after resizing ( $n/m = 4$ )**





## How to delete a key-value pair from a separate-chaining hash table?

- A. Search for key; remove key-value pair from chain.
- B. Compute hash of key; reinsert all other key-value pairs in chain.
- C. Either A or B.
- D. Neither A nor B.



# Symbol table implementations: summary

---

implementation	guarantee			average case			ordered ops?	key interface
	search	insert	delete	search	insert	delete		
<b>sequential search (unordered list)</b>	$n$	$n$	$n$	$n$	$n$	$n$		<code>equals()</code>
<b>binary search (ordered array)</b>	$\log n$	$n$	$n$	$\log n$	$n$	$n$	✓	<code>compareTo()</code>
<b>BST</b>	$n$	$n$	$n$	$\log n$	$\log n$	$\sqrt{n}$	✓	<code>compareTo()</code>
<b>red-black BST</b>	$\log n$	$\log n$	$\log n$	$\log n$	$\log n$	$\log n$	✓	<code>compareTo()</code>
<b>separate chaining</b>	$n$	$n$	$n$	$1^\dagger$	$1^\dagger$	$1^\dagger$		<code>equals()</code> <code>hashCode()</code>

† under uniform hashing assumption



Algorithms

## 3.4 HASH TABLES

---

- ▶ *hash functions*
- ▶ *separate chaining*
- ▶ *linear probing*
- ▶ *context*

ROBERT SEDGEWICK | KEVIN WAYNE

<https://algs4.cs.princeton.edu>

## Linear-probing hash table: insert

---

- Maintain key–value pairs in two parallel arrays, with one key per cell.
- Resolve collisions by probing: search successive cells until either finding the key or an unused cell.

Inserting into a linear-probing hash table.

linear-probing hash table																
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	
keys[]	P	M			A	C		H	L		E			R	X	
vals[]	11	10			9	5		6	12		13			4	8	
put(K, 14)							K									
hash(K) = 7								14								

## Linear-probing hash table: search

---

- Maintain key–value pairs in two parallel arrays, with one key per cell.
- Resolve collisions by probing: search successive cells until either finding the key or an unused cell.

Searching in a linear-probing hash table.

linear-probing hash table																
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	
keys[]	P	M			A	C		H	L	K	E			R	X	
get(K)		get(Z)			K	Z										
hash(K) = 7		hash(Z) = 8														
vals[]	11	10			9	5		6	12	14	13			4	8	



## Linear-probing hash table demo

**Hash.** Map key to integer  $i$  between 0 and  $m - 1$ .

**Insert.** Put at table index  $i$  if free; if not try  $i + 1, i + 2, \dots$

**Search.** Search table index  $i$ ; if occupied but no match, try  $i + 1, i + 2, \dots$

**Note.** Array length  $m$  **must** be greater than number of key-value pairs  $n$ .

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
keys[]	P	M			A	C	S	H	L		E			R	X	

$m = 16$

## Linear-probing symbol table: Java implementation

```
public class LinearProbingHashST<Key, Value>
{
    private int m = 32768;
    private Value[] vals = (Value[]) new Object[m];
    private Key[] keys = (Key[]) new Object[m];

    private int hash(Key key)
    { /* as before */ }

    private void put(Key key, Value val) { /* next slide */ }

    public Value get(Key key)
    {
        for (int i = hash(key); keys[i] != null; i = (i+1) % m)
            if (key.equals(keys[i]))
                return vals[i];
        return null;
    }
}
```

array resizing  
code omitted

## Linear-probing symbol table: Java implementation

---

```
public class LinearProbingHashST<Key, Value>
{
    private int m = 32768;
    private Value[] vals = (Value[]) new Object[m];
    private Key[] keys = (Key[]) new Object[m];


    private int hash(Key key)
    { /* as before */ }

    public Value get(Key key) { /* prev slide */ }

    public void put(Key key, Value val)
    {
        int i;
        for (i = hash(key); keys[i] != null; i = (i+1) % m)
            if (keys[i].equals(key))
                break;
        keys[i] = key;
        vals[i] = val;
    }
}
```



**Under the uniform hashing assumption, where is the next key most likely to be added in this linear-probing hash table (no resizing)?**




- A. Index 7.
- B. Index 14.
- C. Either index 4 or 14.
- D. All open indices are equally likely.

# Clustering

**Cluster.** A contiguous block of keys.

**Observation.** New keys disproportionately likely to hash into big clusters.



# Analysis of linear probing

**Proposition.** Under uniform hashing assumption, the average # of probes in a linear-probing hash table of size  $m$  that contains  $n = \alpha m$  keys is at most


$$\frac{1}{2} \left( 1 + \frac{1}{1 - \alpha} \right)$$

search hit

$$\frac{1}{2} \left( 1 + \frac{1}{(1 - \alpha)^2} \right)$$

search miss / insert

**Pf.** [beyond course scope]



## Parameters.

- $m$  too large  $\Rightarrow$  too many empty array entries.
- $m$  too small  $\Rightarrow$  search time blows up.
- Typical choice:  $\alpha = n / m \sim 1/2.$  ←
  - # probes for search hit is about 3/2
  - # probes for search miss is about 5/2

# Resizing in a linear-probing hash table

**Goal.** Average length of list  $n / m \leq \frac{1}{2}$ .

- Double length of array  $m$  when  $n / m \geq \frac{1}{2}$ .
- Halve length of array  $m$  when  $n / m \leq \frac{1}{8}$ .
- Need to rehash all keys when resizing.

**before resizing**

	0	1	2	3	4	5	6	7
keys[]		E	S			R	A	
vals[]		1	0			3	2	

**after resizing**

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
keys[]				A		S			E				R			
vals[]					2		0			1				3		



## How to delete a key-value pair from a linear-probing hash table?

A. Search for key; remove key-value pair from arrays.

B. Search for key; remove key-value pair from arrays.

Shift all keys in **cluster** after deleted key 1 position to left.

C. Either A and B.

D. Neither A nor B.

cluster after deleted key																
before deleting S																
keys[]	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
vals[]	P	M			A	C	S	H	L		E			R	X	
	10	9			8	4	0	5	11		12			3	7	

# ST implementations: summary

---

implementation	guarantee			average case			ordered ops?	key interface
	search	insert	delete	search	insert	delete		
sequential search (unordered list)	$n$	$n$	$n$	$n$	$n$	$n$		<code>equals()</code>
binary search (ordered array)	$\log n$	$n$	$n$	$\log n$	$n$	$n$	✓	<code>compareTo()</code>
BST	$n$	$n$	$n$	$\log n$	$\log n$	$\sqrt{n}$	✓	<code>compareTo()</code>
red-black BST	$\log n$	$\log n$	$\log n$	$\log n$	$\log n$	$\log n$	✓	<code>compareTo()</code>
separate chaining	$n$	$n$	$n$	$1^\dagger$	$1^\dagger$	$1^\dagger$		<code>equals()</code> <code>hashCode()</code>
linear probing	$n$	$n$	$n$	$1^\dagger$	$1^\dagger$	$1^\dagger$		<code>equals()</code> <code>hashCode()</code>

† under uniform hashing assumption

# 3-SUM (REVISITED)



**3-SUM.** Given  $n$  distinct integers, find three such that  $a + b + c = 0$ .

**Goal.**  $\Theta(n^2)$  expected time;  $\Theta(n)$  extra space.



Algorithms

## 3.4 HASH TABLES

---

- ▶ *hash functions*
- ▶ *separate chaining*
- ▶ *linear probing*
- ▶ ***context***

ROBERT SEDGEWICK | KEVIN WAYNE

<https://algs4.cs.princeton.edu>


## War story: algorithmic complexity attacks

---

Q. Is the uniform hashing assumption important in practice?

A1. Yes: aircraft control, nuclear reactor, pacemaker, HFT, ...

A2. Yes: denial-of-service (DoS) attacks.



Real-world exploits. [Crosby–Wallach 2003]

- Linux 2.4.20 kernel: save files with carefully chosen names.
- Bro server: send carefully chosen packets to DoS the server, using less bandwidth than a dial-up modem.

# War story: algorithmic complexity attacks

## A Java bug report.


Jan Lieskovsky 2011-11-01 14:13:47 UTC	Description
Julian Wälde and Alexander Klink reported that the <code>String.hashCode()</code> hash function is not sufficiently collision resistant. <code>hashCode()</code> value is used in the implementations of <code>HashMap</code> and <code>Hashtable</code> classes:	
<a href="http://docs.oracle.com/javase/6/docs/api/java/util/HashMap.html">http://docs.oracle.com/javase/6/docs/api/java/util/HashMap.html</a> <a href="http://docs.oracle.com/javase/6/docs/api/java/util/Hashtable.html">http://docs.oracle.com/javase/6/docs/api/java/util/Hashtable.html</a>	
A specially-crafted set of keys could trigger hash function collisions, which can degrade performance of <code>HashMap</code> or <code>Hashtable</code> by changing hash table operations complexity from an expected/average $O(1)$ to the worst case $O(n)$ . Reporters were able to find colliding strings efficiently using equivalent substrings and meet in the middle techniques.	
This problem can be used to start a denial of service attack against Java applications that use untrusted inputs as <code>HashMap</code> or <code>Hashtable</code> keys. An example of such application is web application server (such as tomcat, see <a href="#">bug #750521</a> ) that may fill hash tables with data from HTTP request (such as GET or POST parameters). A remote attack could use that to make JVM use excessive amount of CPU time by sending a POST request with large amount of parameters which hash to the same value.	
This problem is similar to the issue that was previously reported for and fixed in e.g. perl: <a href="http://www.cs.rice.edu/~scrosby/hash/CrosbyWallach_UseNixSec2003.pdf">http://www.cs.rice.edu/~scrosby/hash/CrosbyWallach_UseNixSec2003.pdf</a>	

[https://bugzilla.redhat.com/show\\_bug.cgi?id=750533](https://bugzilla.redhat.com/show_bug.cgi?id=750533)

# Hashing: file verification

When downloading a file from the web:

- Vendor publishes hash of file.
- Client checks whether hash of downloaded file matches.
- If mismatch, file corrupted. ← (e.g., error in transmission or infected by virus)



```
~/Desktop> sha256sum ideaIC-2019.3.3.dmg  
c62ed2df891ccbb40d890e8a0074781801f086a3091a4a2a592a96afaba31270
```

# Hashing: cryptographic applications

---

One-way hash function. “Hard” to find a key that will hash to a desired value  
(or two keys that hash to same value).

Ex. MD5, SHA-1, SHA-256, SHA-512, Whirlpool, ....

known to be insecure




Applications. File verification, digital signatures, cryptocurrencies,  
password authentication, blockchain, Git commit identifiers, ....

# Separate chaining vs. linear probing

## Separate chaining.

- Performance degrades gracefully.
- Clustering less sensitive to poorly-designed hash function.



## Linear probing.

- Less memory.
- Better cache performance.
- More probes because of clustering.

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
keys[]	P	M			A	C	S	H	L		E				R	X
vals[]	10	9			8	4	0	5	11		12				3	7

# Hashing: variations on the theme

---

Many improved versions have been studied.

## Two-probe hashing. [ separate-chaining variant ]


- Hash to two positions, insert key in shorter of the two chains.
- Reduces expected length of the longest chain to  $\Theta(\log \log n)$ .

## Double hashing. [ linear-probing variant ]

- Resolve collisions by probing, but skip a variable amount instead of +1.
- Effectively eliminates clustering.
- Can allow table to become nearly full.
- More difficult to implement delete.

## Cuckoo hashing. [ linear-probing variant ]

- Hash key to two positions; insert key into either position; if occupied, reinsert displaced key into its alternative position (and recur).
- $\Theta(1)$  time for search in worst case.




# Hash tables vs. balanced search trees

---


## Hash tables.

- Simpler to code.
- Typically faster in practice.
- No effective alternative for unordered keys.



## Balanced search trees.

- Stronger performance guarantee.
- Support for ordered ST operations.
- Easier to implement `compareTo()` than `hashCode()`.



## Java includes both.

- BSTs: `java.util.TreeMap`, `java.util.TreeSet`. ← red-black BST
- Hash tables: `java.util.HashMap`, `java.util.HashSet`, `java.util.IdentityHashMap`.

↑  
separate chaining  
(Java 8: if chain gets too long,  
use red-black BST for chain)

↑  
linear probing

© Copyright 2021 Robert Sedgewick and Kevin Wayne