



<https://algs4.cs.princeton.edu>

## 5.1 STRING SORTS

---

- ▶ *strings in Java*
- ▶ *key-indexed counting*
- ▶ *LSD radix sort*
- ▶ *MSD radix sort*
- ▶ *3-way radix quicksort*
- ▶ *suffix arrays*



## 5.1 STRING SORTS

---

- ▶ *strings in Java*
- ▶ *key-indexed counting*
- ▶ *LSD radix sort*
- ▶ *MSD radix sort*
- ▶ *3-way radix quicksort*
- ▶ *suffix arrays*

# String processing

---

String. Sequence of characters.

Important fundamental abstraction.

- Programming systems (e.g., Java programs).
- Communication systems (e.g., email).
- Information processing.
- Genomic sequences.
- ...

*“The digital information that underlies biochemistry, cell biology, and development can be represented by a simple string of G’s, A’s, T’s and C’s. This string is the root data structure of an organism’s biology. ” — M. V. Olson*



# The char data type

C **char** data type. Typically an 8-bit integer.

- Supports 7-bit ASCII.
- Represents only  $2^8 = 256$  characters.

0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F
NUL	SOH	STX	ETX	EOT	ENQ	ACK	BEL	BS	HT	LF	VT	FF	CR	SO	SI
DLE	DC1	DC2	DC3	DC4	NAK	SYN	ETB	CAN	EM	SUB	ESC	FS	GS	RS	US
SP	!	"	#	\$	%	&	'	(	)	*	+	,	-	.	/
0	1	2	3	4	5	6	7	8	9	:	;	<	=	>	?
@	A	B	C	D	E	F	G	H	I	J	K	L	M	N	0
P	Q	R	S	T	U	V	W	X	Y	Z	[	\	]	^	_
`	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o
p	q	r	s	t	u	v	w	x	y	z	{		}	~	DEL

all  $2^7 = 128$  ASCII characters

A á ð 💩

U+0041 U+00E1 U+2202 U+1F4A9

some Unicode characters

can use as an index into an array



Java **char** data type. A 16-bit unsigned integer (between 0 and  $2^{16} = 65,536$ ).

- Supports 16-bit Unicode 1.0.1. ← 7,161 characters
- Supports 21-bit Unicode 10.0.0 (via UTF-8). ← 136,755 characters and emoji


U+1F496



# The String data type (in Java)

String data type. Immutable sequence of characters.

Java representation. A fixed-length char[] array.



operation	description	Java	running time
<b>length</b>	<i>number of characters</i>	<code>s.length()</code>	1
<b>indexing</b>	<i>character at index <math>i</math></i>	<code>s.charAt(i)</code>	1
<b>concatenation</b>	<i>concatenate one string to the end of the other</i>	<code>s + t</code>	$m + n$
:	:		

# String performance trap

Q. How to build a long string, one character at a time?

```
public static String reverse(String s)
{
    String reverse = "";
    for (int i = s.length() - 1; i >= 0; i--)
        reverse += s.charAt(i);
    return reverse;
}
```

quadratic time  
 $(1 + 2 + 3 + \dots + n)$

StringBuilder data type. Mutable sequence of characters.

Java representation. A resizing char[] array.

or equivalently,  
new StringBuilder(s).reverse().toString()

```
public static String reverse(String s)
{
    StringBuilder reverse = new StringBuilder();
    for (int i = s.length() - 1; i >= 0; i--)
        reverse.append(s.charAt(i));
    return reverse.toString();
}
```

linear time  
 $n + (1 + 2 + 4 + 8 + \dots + n)$

# THE STRING DATA TYPE: IMMUTABILITY



Q. Why are Java strings immutable?

# Alphabets

---

**Digital key.** Sequence of digits over a given alphabet.

**Radix.** Number of digits  $R$  in alphabet.

name	$R()$	$\lg R()$	characters
BINARY	2	1	01
OCTAL	8	3	01234567
DECIMAL	10	4	0123456789
HEXADECIMAL	16	4	0123456789ABCDEF
DNA	4	2	ACTG
LOWERCASE	26	5	abcdefghijklmnopqrstuvwxyz
UPPERCASE	26	5	ABCDEFGHIJKLMNOPQRSTUVWXYZ
PROTEIN	20	5	ACDEFGHIJKLMNOPQRSTUVWXYZ
BASE64	64	6	ABCDEFGHIJKLMNOPQRSTUVWXYZabcdefghijklmnopqrstuvwxyz0123456789+/
ASCII	128	7	<i>ASCII characters</i>
EXTENDED_ASCII	256	8	<i>extended ASCII characters</i>
UNICODE16	65536	16	<i>Unicode characters</i>

**Bottom line.** We assume ASCII strings; extends to 64-bit integers and other digital keys.



## 5.1 STRING SORTS

---

- ▶ *strings in Java*
- ▶ **key-indexed counting**
- ▶ *LSD radix sort*
- ▶ *MSD radix sort*
- ▶ *3-way radix quicksort*
- ▶ *suffix arrays*

# Review: summary of the performance of sorting algorithms

Frequency of calls to `compareTo()`.

algorithm	guarantee	random	extra space	stable?	operations on keys
insertion sort	$\frac{1}{2} n^2$	$\frac{1}{4} n^2$	$\Theta(1)$	✓	<code>compareTo()</code>
mergesort	$n \log_2 n$	$n \log_2 n$	$\Theta(n)$	✓	<code>compareTo()</code>
quicksort	$1.39 n \log_2 n$ *	$1.39 n \log_2 n$ *	$\Theta(\log n)$ *		<code>compareTo()</code>
heapsort	$2 n \log_2 n$	$2 n \log_2 n$	$\Theta(1)$		<code>compareTo()</code>

\* probabilistic

Sorting lower bound. In the worst case, any compare-based sorting algorithm

makes  $\Omega(n \log n)$  compares. ← compareTo() not constant time for string keys

Q. Can we sort strings faster (despite lower bound)?

A. Yes, by exploiting access to individual characters. ← use characters to make  
R-way decisions  
(instead of binary decisions)

## Key-indexed counting: assumptions about keys

**Assumption.** Each key is an integer between 0 and  $R - 1$ .

**Implication.** Can use key as an array index.

### Applications.

- Sort string by first letter.
- Sort playing cards by suit.
- Sort phone numbers by area code.
- Sort class roster by section number.
- Use as a subroutine in string sorting algorithm.

**Remark.** Keys typically have associated data  $\Rightarrow$   
can't simply count keys of each value.

input		sorted result (by section)	
<i>name</i>	<i>section</i>		
Anderson	2	Harris	1
Brown	3	Martin	1
Davis	3	Moore	1
Garcia	4	Anderson	2
Harris	1	Martinez	2
Jackson	3	Miller	2
Johnson	4	Robinson	2
Jones	3	White	2
Martin	1	Brown	3
Martinez	2	Davis	3
Miller	2	Jackson	3
Moore	1	Jones	3
Robinson	2	Taylor	3
Smith	4	Williams	3
Taylor	3	Garcia	4
Thomas	4	Johnson	4
Thompson	4	Smith	4
White	2	Thomas	4
Williams	3	Thompson	4
Wilson	4	Wilson	4

↑  
*keys are  
small integers*



# Key-indexed counting demo

**Goal.** Sort an array  $a[]$  of  $n$  characters between 0 and  $R - 1$ .

- Count frequencies of each letter using key as index.
- Compute frequency cumulates which specify destinations.
- Distribute items to auxiliary array using cumulative frequencies.
- Copy back into original array.

$R = 6$

```

int n = a.length;
int[] count = new int[R+1];

for (int i = 0; i < n; i++)
    count[a[i]+1]++;

for (int r = 0; r < R; r++)
    count[r+1] += count[r];

for (int i = 0; i < n; i++)
    aux[count[a[i]]++] = a[i];

for (int i = 0; i < n; i++)
    a[i] = aux[i];

```

i	a[i]
0	d
1	a
2	c
3	f
4	f
5	b
6	d
7	b
8	f
9	b
10	e
11	a

use a for 0  
 b for 1  
 c for 2  
 d for 3  
 e for 4  
 f for 5



# Key-indexed counting demo

**Goal.** Sort an array  $a[]$  of  $n$  characters between 0 and  $R - 1$ .

- Compute character frequencies.
- Compute cumulative frequencies.
- Distribute items to auxiliary array using cumulative frequencies.
- Copy back into original array.

```

int n = a.length;
int[] count = new int[R+1];

for (int i = 0; i < n; i++)
    count[a[i]+1]++;
}

for (int r = 0; r < R; r++)
    count[r+1] += count[r];

for (int i = 0; i < n; i++)
    aux[count[a[i]]++] = a[i];

for (int i = 0; i < n; i++)
    a[i] = aux[i];

```

count  
frequencies

i	a[i]	offset by 1 [stay tuned]
0	d	
1	a	
2	c	
3	f	
4	f	
5	b	
6	d	
7	b	
8	f	
9	b	
10	e	
11	a	

offset by 1  
[stay tuned]

a	0
b	2
c	3
d	1
e	2
f	1
-	3

r count[r]



# Key-indexed counting demo

**Goal.** Sort an array  $a[]$  of  $n$  characters between 0 and  $R - 1$ .

- Compute character frequencies.
- Compute cumulative frequencies.
- Distribute items to auxiliary array using cumulative frequencies.
- Copy back into original array.

```

int n = a.length;
int[] count = new int[R+1];

for (int i = 0; i < n; i++)
    count[a[i]+1]++;

for (int r = 0; r < R; r++)
    count[r+1] += count[r];

for (int i = 0; i < n; i++)
    aux[count[a[i]]++] = a[i];

for (int i = 0; i < n; i++)
    a[i] = aux[i];

```

compute  
cumulates

i	a[i]	r	count[r]
0	d	a	0
1	a	b	2
2	c	c	5
3	f	d	6
4	f	e	8
5	b	f	9
6	d		12
7	b		
8	f		
9	b		
10	e		
11	a		

6 keys < d, 8 keys < e  
so d's go in a[6] and a[7]



# Key-indexed counting demo

**Goal.** Sort an array  $a[]$  of  $n$  characters between 0 and  $R - 1$ .

- Compute character frequencies.
- Compute cumulative frequencies.
- Distribute items to auxiliary array using cumulative frequencies.
- Copy back into original array.

```

int n = a.length;
int[] count = new int[R+1];

for (int i = 0; i < n; i++)
    count[a[i]+1]++;
}

for (int r = 0; r < R; r++)
    count[r+1] += count[r];

for (int i = 0; i < n; i++)
    aux[count[a[i]]++] = a[i];

for (int i = 0; i < n; i++)
    a[i] = aux[i];

```

move  
items

i	a[i]	i	aux[i]
0	d	0	a
1	a	1	a
2	c	2	b
3	f	3	b
4	f	4	b
5	b	5	c
6	d	6	d
7	b	7	d
8	f	8	e
9	b	9	f
10	e	10	f
11	a	11	f

r count[r]

a	2
b	5
c	6
d	8
e	9
f	12
-	12



# Key-indexed counting demo

**Goal.** Sort an array  $a[]$  of  $n$  characters between 0 and  $R - 1$ .

- Compute character frequencies.
- Compute cumulative frequencies.
- Distribute items to auxiliary array using cumulative frequencies.
- Copy back into original array.

```

int n = a.length;
int[] count = new int[R+1];

for (int i = 0; i < n; i++)
    count[a[i]+1]++;
}

for (int r = 0; r < R; r++)
    count[r+1] += count[r];

for (int i = 0; i < n; i++)
    aux[count[a[i]]++] = a[i];

for (int i = 0; i < n; i++)
    a[i] = aux[i];

```


copy  
back

i	a[i]	i	aux[i]
0	a	0	a
1	a	1	a
2	b	2	b
3	b	3	b
4	b	4	b
5	c	5	c
6	d	6	d
7	d	7	d
8	e	8	e
9	f	9	f
-		10	f
10	f	11	f
11	f		



Which of the following are properties of key-indexed counting?

- A.  $\Theta(n + R)$  time.
- B.  $\Theta(n + R)$  extra space.
- C. Stable.
- D. All of the above.






## 5.1 STRING SORTS

---

- ▶ *strings in Java*
- ▶ *key-indexed counting*
- ▶ **LSD radix sort**
- ▶ *MSD radix sort*
- ▶ *3-way radix quicksort*
- ▶ *suffix arrays*

# Least-significant-digit-first (LSD) radix sort

- Consider characters from **right to left**.
- Stably sort using character  $d$  as the key (using key-indexed counting).




## LSD string sort: correctness proof

Proposition. LSD sorts any array of  $n$  strings, each of length  $w$ , in  $\Theta(w(n + R))$  time.

Pf of correctness. [ by induction on # passes  $i$  ]

- Inductive hypothesis: after pass  $i$ , strings are sorted by last  $i$  characters.
- After pass  $i + 1$ , strings are sorted by last  $i + 1$  last characters because...
  - if two strings differ on sort key, key-indexed counting puts them in proper relative order
  - if two strings agree on sort key, stability of key-indexed counting keeps them in proper relative order



Proposition. LSD sort is stable.

Pf. Key-indexed counting is stable.

## LSD string sort (for fixed-length strings): Java implementation

```
public class LSD
{
    public static void sort(String[] a, int w)
    {
        int R = 256; ← radix R
        int n = a.length;
        String[] aux = new String[n];

        for (int d = w-1; d >= 0; d--) ←
        {
            int[] count = new int[R+1];      key-indexed counting
            for (int i = 0; i < n; i++)       (using character d)
                count[a[i].charAt(d) + 1]++;
            for (int r = 0; r < R; r++)
                count[r+1] += count[r];
            for (int i = 0; i < n; i++)
                aux[count[a[i].charAt(d)]++] = a[i];
            for (int i = 0; i < n; i++)
                a[i] = aux[i];
        }
    }
}
```

fixed-length  $w$  strings

do key-indexed counting  
for each digit from right to left

key-indexed counting  
(using character  $d$ )

# Summary of the performance of sorting algorithms

---

Frequency of calls to `compareTo()` and `charAt()`.

algorithm	guarantee	random	extra space	stable?	operations on keys
insertion sort	$\frac{1}{2} n^2$	$\frac{1}{4} n^2$	$\Theta(1)$	✓	<code>compareTo()</code>
mergesort	$n \log_2 n$	$n \log_2 n$	$\Theta(n)$	✓	<code>compareTo()</code>
quicksort	$1.39 n \log_2 n$ *	$1.39 n \log_2 n$ *	$\Theta(\log n)$ *		<code>compareTo()</code>
heapsort	$2 n \log_2 n$	$2 n \log_2 n$	$\Theta(1)$		<code>compareTo()</code>
LSD sort †	$2 w n$	$2 w n$	$\Theta(n + R)$	✓	<code>charAt()</code>

\* probabilistic

† fixed-length  $w$  keys



1 call to `compareTo()` can involve as many as  $2w$  calls to `charAt()`




Google CEO Eric Schmidt interviews Barack Obama in November 2007



Which algorithm below is fastest for sorting 1 million 32-bit integers?

- A. Insertion sort.
- B. Mergesort.
- C. Quicksort.
- D. LSD sort.





## 5.1 STRING SORTS

---

- ▶ *strings in Java*
- ▶ *key-indexed counting*
- ▶ *LSD radix sort*
- ▶ ***MSD radix sort***
- ▶ *3-way radix quicksort*
- ▶ *suffix arrays*

## Reverse LSD

- Consider characters from **left to right**.
- Stably sort using character  $d$  as the key (using key-indexed counting).

The diagram illustrates the Reverse LSD sorting process through three stages of sorting:

- sort key ( $d = 0$ )**: The first stage sorts by the leftmost character. Characters 'b' and 'd' are highlighted in red. Arrows point from the second column of the first stage to the first column of the second stage.
- sort key ( $d = 1$ )**: The second stage sorts by the middle character. Characters 'b' and 'd' are highlighted in red. Arrows point from the third column of the first stage to the second column of the second stage.
- sort key ( $d = 2$ )**: The third stage sorts by the rightmost character. Characters 'c', 'd', 'e', and 'f' are highlighted in grey. Arrows point from the first column of the second stage to the third column of the third stage.

The final sorted state is shown below the third stage:

0	c	a	b
1	d	a	b
2	e	b	b
3	f	a	d
4	a	c	e
5	b	d	d
6	c	e	e
7	d	f	e
8	e	a	d
9	f	b	d
10	a	c	e
11	b	d	e

**not sorted!**


# Most-significant-digit-first (MSD) radix sort

## Overview.

- Partition array into  $R$  subarrays according to first character. ← use key-indexed counting
- Recursively sort all strings that start with each character. ← key-indexed counts delineate subarray boundaries  
(excluding the first characters in subsequent sorts)

0	d	a	b
1	a	d	d
2	c	a	b
3	f	a	d
4	f	e	e
5	b	a	d
6	d	a	d
7	b	e	e
8	f	e	d
9	b	e	d
10	e	b	b
11	a	c	e

sort key ( $d = 0$ )



# MSD string sort (for fixed-length strings): Java implementation

```
public static void sort(String[] a, int w) ← fixed-length  $w$  strings
{
    aux = new String[a.length]; ←
    sort(a, aux, w, 0, a.length - 1, 0); ←
}

private static void sort(String[] a, String[] aux, int w, int lo, int hi, int d) ←
{
    if (hi <= lo || d == w) return; ← subarrays of length 0 or 1; or all  $w$  characters match
    int[] count = new int[R+1]; ← key-indexed counting
    for (int i = lo; i <= hi; i++) (using character  $d$ )
        count[a[i].charAt(d) + 1]++;
    for (int r = 0; r < R; r++)
        count[r+1] += count[r];
    for (int i = lo; i <= hi; i++)
        aux[count[a[i].charAt(d)]++] = a[i];
    for (int i = lo; i <= hi; i++)
        a[i] = aux[i - lo];
}

sort(a, aux, w, lo, lo + count[0] - 1, d+1); ← sort  $R$  subarrays recursively
for (int r = 1; r < R; r++)
    sort(a, aux, w, lo + count[r-1], lo + count[r] - 1, d+1); ←
}

at this place in code, count[r] = number of keys ≤ r
```

# Variable-length strings

Useful trick. Treat strings as if they had an extra char at end (smaller than any char).

0	s	e	a	-1
1	s	e	a	s
2	s	e	l	l
3	s	h	e	-1
4	s	h	e	-1
5	s	h	e	l
6	s	h	o	r

why smaller?

"she" before "shells"

```
private static int charAt(String s, int d)
{
    if (d < s.length()) return s.charAt(d);
    else return -1;
}
```

C strings. Terminated with null character ('\0')  $\Rightarrow$  no extra work needed.



For which family of inputs is MSD sort likely to be faster than LSD sort?

- A. Random strings.
- B. All equal strings.
- C. Both A and B.
- D. Neither A nor B.

random	all equal
1 E I 0 4 0 2	1 D N B 3 7 7
1 H Y L 4 9 0	1 D N B 3 7 7
1 R O Z 5 7 2	1 D N B 3 7 7
2 H X E 7 3 4	1 D N B 3 7 7
2 I Y E 2 3 0	1 D N B 3 7 7
2 X O R 8 4 6	1 D N B 3 7 7
3 C D B 5 7 3	1 D N B 3 7 7
3 C V P 7 2 0	1 D N B 3 7 7
3 I G J 3 1 9	1 D N B 3 7 7
3 K N A 3 8 2	1 D N B 3 7 7
3 T A V 8 7 9	1 D N B 3 7 7
4 C Q P 7 8 1	1 D N B 3 7 7
4 Q G I 2 3 4	1 D N B 3 7 7
4 Y H V 2 2 9	1 D N B 3 7 7

## MSD string sort: performance

---

**Observation.** MSD examines just enough character to sort the keys.

**Proposition.** For random strings, MSD examines  $\Theta(n \log_R n)$  characters.

**Remark.** This can be sublinear in the input size  $\Theta(n w)$ . ← compareTo() based sorts can also be sublinear

**Proposition.** In the worst case, MSD requires  $\Theta(n + wR)$  extra space.

random	all equal
1 E I 0 4 0 2	1 D N B 3 7 7
1 H Y L 4 9 0	1 D N B 3 7 7
1 R O Z 5 7 2	1 D N B 3 7 7
2 H X E 7 3 4	1 D N B 3 7 7
2 I Y E 2 3 0	1 D N B 3 7 7
2 X O R 8 4 6	1 D N B 3 7 7
3 C D B 5 7 3	1 D N B 3 7 7
3 C V P 7 2 0	1 D N B 3 7 7
3 I G J 3 1 9	1 D N B 3 7 7
3 K N A 3 8 2	1 D N B 3 7 7
3 T A V 8 7 9	1 D N B 3 7 7
4 C Q P 7 8 1	1 D N B 3 7 7
4 Q G I 2 3 4	1 D N B 3 7 7
4 Y H V 2 2 9	1 D N B 3 7 7

# Summary of the performance of sorting algorithms

Frequency of `compareTo()` and `charAt()` operations.

algorithm	guarantee	random	extra space	stable?	operations on keys
insertion sort	$\frac{1}{2} n^2$	$\frac{1}{4} n^2$	$\Theta(1)$	✓	<code>compareTo()</code>
mergesort	$n \log_2 n$	$n \log_2 n$	$\Theta(n)$	✓	<code>compareTo()</code>
quicksort	$1.39 n \log_2 n$ *	$1.39 n \log_2 n$ *	$\Theta(\log n)$ *		<code>compareTo()</code>
heapsort	$2 n \log_2 n$	$2 n \log_2 n$	$\Theta(1)$		<code>compareTo()</code>
LSD sort †	$2 w n$	$2 w n$	$\Theta(n + R)$	✓	<code>charAt()</code>
MSD sort ‡	$2 w n$	$n \log_R n$	$\Theta(n + D R)$	✓	<code>charAt()</code>

running time can be  $\Theta(wnR)$   
( $n/2$  pairs of duplicate keys)

$D$  = function-call stack depth  
(length of longest common prefix)

\* probabilistic

† fixed-length  $w$  keys

‡ average-length  $w$  keys

# Engineering a radix sort (American flag sort)

## Optimization 0. Cutoff to insertion sort.

- MSD is much too slow for small subarrays.
- Essential for performance.

## Optimization 1. Replace recursion with explicit stack.

- Push subarrays to be sorted onto stack.
- One count[] array now suffices.

## Optimization 2. Do $R$ -way partitioning in place.

- Eliminates aux[] array.
- Sacrifices stability.



American national flag problem



Dutch national flag problem

*Engineering Radix Sort*

Peter M. McIlroy and Keith Bostic  
University of California at Berkeley;  
and M. Douglas McIlroy  
AT&T Bell Laboratories

---

ABSTRACT: Radix sorting methods have excellent asymptotic performance on string data, for which comparison is not a unit-time operation. Attractive for use in large byte-addressable memories, these methods have nevertheless long been eclipsed by more easily programmed algorithms. Three ways to sort strings by bytes left to right—a stable list sort, a stable two-array sort, and an in-place “American flag” sort—are illustrated with practical C programs. For heavy-duty sorting, all three perform comparably, usually running at least twice as fast as a good quicksort. We recommend American flag sort for general use.



## 5.1 STRING SORTS


---

- ▶ *strings in Java*
- ▶ *key-indexed counting*
- ▶ *LSD radix sort*
- ▶ *MSD radix sort*
- ▶ ***3-way radix quicksort***
- ▶ *suffix arrays*


# 3-way string quicksort

## Overview.

- Partition array into 3 subarrays according to first character of pivot. ←———— use Dijkstra 3-way partitioning algorithm
- Recursively sort 3 subarrays. ←———— exclude first character when sorting middle subarray (since known to be equal)



# 3-way string quicksort: trace of recursive calls



## Trace of first few recursive calls for 3-way string quicksort (subarrays of length 1 not shown)

## 3-way string quicksort: Java implementation

```
private static void sort(String[] a)
{  sort(a, 0, a.length - 1, 0); }
```

```
private static void sort(String[] a, int lo, int hi, int d) ←
{
    if (hi <= lo) return; ← subarrays of length 0 or 1
    int pivot = charAt(a[lo], d);
```

```
int lt = lo, gt = hi;
int i = lo + 1;
while (i <= gt)
{
    int c = charAt(a[i], d);
    if      (c < pivot) exch(a, lt++, i++);
    else if (c > pivot) exch(a, i, gt--);
    else
        i++;
}
```

Dijkstra 3-way partitioning  
(using character at index  $d$ )

sort  $a[lo..hi]$  assuming first  $d$   
characters are equal

```
sort(a, lo, lt-1, d);          sort 3 subarrays recursively
if (pivot != -1) sort(a, lt, gt, d+1);
sort(a, gt+1, hi, d);
```

```
}
```



# 3-way string quicksort vs. competitors

## 3-way string quicksort vs. MSD sort.

- In-place; short inner loop; cache-friendly.
- Not stable.

## 3-way string quicksort vs. standard quicksort.

- Typically uses  $\sim 2n \ln n$  character compares (instead of  $\sim 2n \ln n$  string compares).
- Faster for keys with long common prefixes (and this is a common case!)



library of Congress call numbers

Bottom line. 3-way string quicksort is often the method of choice for sorting strings.

# Summary of the performance of sorting algorithms

---

Frequency of `compareTo()` and `charAt()` operations.

algorithm	guarantee	random	extra space	stable?	operations on keys
<b>insertion sort</b>	$\frac{1}{2} n^2$	$\frac{1}{4} n^2$	$\Theta(1)$	✓	<code>compareTo()</code>
<b>mergesort</b>	$n \log_2 n$	$n \log_2 n$	$\Theta(n)$	✓	<code>compareTo()</code>
<b>quicksort</b>	$1.39 n \log_2 n$ *	$1.39 n \log_2 n$ *	$\Theta(\log n)$ *		<code>compareTo()</code>
<b>heapsort</b>	$2 n \log_2 n$	$2 n \log_2 n$	$\Theta(1)$		<code>compareTo()</code>
<b>LSD sort</b> †	$2 w n$	$2 w n$	$\Theta(n + R)$	✓	<code>charAt()</code>
<b>MSD sort</b> ‡	$2 w n$	$n \log_R n$	$\Theta(n + DR)$	✓	<code>charAt()</code>
<b>3-way string quicksort</b>	$1.39 w n \log_2 R$ *	$1.39 n \log_2 n$ *	$\Theta(\log n + w)$ *		<code>charAt()</code>

\* probabilistic

† fixed-length  $w$  keys

‡ average-length  $w$  keys



## 5.1 STRING SORTS

---

- ▶ *strings in Java*
- ▶ *key-indexed counting*
- ▶ *LSD radix sort*
- ▶ *MSD radix sort*
- ▶ *3-way radix quicksort*
- ▶ ***suffix arrays***

## Keyword-in-context search

Given a text of  $n$  characters, preprocess it to enable fast substring search  
(find all occurrences of query string context).

number of characters of surrounding context

```
~/Desktop/51radix> java KWIC tale.txt 15
search
o st giless to search for contraband
her unavailing search for your fathe
le and gone in search of her husband
t provinces in search of impoverishe
dispersing in search of other carri
n that bed and search the straw hold

the epoch
ishness it was the epoch of belief it w
belief it was the epoch of incredulity
```

```
~/Desktop/51radix> more tale.txt
it was the best of times
it was the worst of times
it was the age of wisdom
it was the age of foolishness
it was the epoch of belief
it was the epoch of incredulity
it was the season of light
it was the season of darkness
it was the spring of hope
it was the winter of despair
...
```

**Applications.** Linguistics, databases, web search, word processing, ....

# Suffix sort

input string

i t w a s b e s t i t w a s s w  
0 1 2 3 4 5 6 7 8 9 10 11 12 13 14

form suffixes

i t w a s b e s t i t w a s s w  
t w a s b e s t i t w a s s w  
w a s b e s t i t w a s s w  
a s b e s t i t w a s s w  
s b e s t i t w a s s w  
b e s t i t w a s s w  
e s t i t w a s s w  
s t i t w a s s w  
t i t w a s s w  
i t w a s s w  
t w a s s w  
w a s s w  
a s s w  
s s w  
s w

sort suffixes to bring query strings together

3 a s b e s t i t w a s s w  
12 a s s w  
5 b e s t i t w a s s w  
6 e s t i t w a s s w  
0 i t w a s s b e s t i t w a s s w  
9 i t w a s s w  
4 s b e s t i t w a s s w  
7 s t i t w a s s w  
13 s w  
8 t i t w a s s w  
1 t w a s s b e s t i t w a s s w  
10 t w a s s w  
14 w  
2 w a s s b e s t i t w a s s w  
11 w a s s w

array of suffix indices  
(in sorted order)

# Keyword-in-context search: suffix-sorting solution

- Preprocess: **suffix sort** the text.
- Query: **binary search** for query; scan until mismatch.

**KWIC search for “search” in Tale of Two Cities**

	:
632698	s e a l e d _ m y _ l e t t e r _ a n d _ ...
713727	s e a m s t r e s s _ i s _ l i f t e d _ ...
660598	s e a m s t r e s s _ o f _ t w e n t y _ ...
67610	s e a m s t r e s s _ w h o _ w a s _ w i ...
→ 4430	s e a r c h _ f o r _ c o n t r a b a n d ...
42705	s e a r c h _ f o r _ y o u r _ f a t h e ...
499797	s e a r c h _ o f _ h e r _ h u s b a n d ...
182045	s e a r c h _ o f _ i m p o v e r i s h e ...
143399	s e a r c h _ o f _ o t h e r _ c a r r i ...
411801	s e a r c h _ t h e _ s t r a w _ h o l d ...
158410	s e a r e d _ m a r k i n g _ a b o u t _ ...
691536	s e a s e _ a n d _ m a d a m e _ d e f a r ...
536569	s e a s e _ a _ t e r r i b l e _ p a s s ...
484763	s e a s e _ t h a t _ h a d _ b r o u g h ...
	:



## How much memory as a function of n?

```
String[] suffixes = new String[n];
for (int i = 0; i < n; i++)
    suffixes[i] = s.substring(i, n);

Arrays.sort(suffixes);
```



3<sup>rd</sup> printing (2012)


- A.  $\Theta(1)$
- B.  $\Theta(n)$
- C.  $\Theta(n \log n)$
- D.  $\Theta(n^2)$

# War story

Q. How to efficiently form (and sort) the  $n$  suffixes?

```
String[] suffixes = new String[n];
for (int i = 0; i < n; i++)
    suffixes[i] = s.substring(i, n);

Arrays.sort(suffixes);
```



3<sup>rd</sup> printing (2012)

input file	characters	Java 7u5	Java 7u6
amendments.txt	18 K	0.25 sec	2.0 sec
aesop.txt	192 K	1.0 sec	<i>out of memory</i>
mobydick.txt	1.2 M	7.6 sec	<i>out of memory</i>
chromosome11.txt	7.1 M	61 sec	<i>out of memory</i>

# The String data type: Java 7u6 implementation

---

```
public final class String implements Comparable<String>
{
    private char[] value; // characters
    private int hash;    // cache of hashCode()
    ...
}
```

**String s = "Hello, World";**

value[]	H	E	L	L	0	,		W	0	R	L	D
	0	1	2	3	4	5	6	7	8	9	10	11

**String t = s.substring(7, 12);**

(linear extra memory)

value[]	W	0	R	L	D
	0	1	2	3	4


# The String data type: Java 7u5 implementation

```
public final class String implements Comparable<String>
{
    private char[] value;      // characters
    private int offset;        // index of first char in array
    private int length;        // length of string
    private int hash;          // cache of hashCode()

    ...
}
```

**String s = "Hello, World";**

length = 12



**String t = s.substring(7, 12);**

(constant extra memory)

length = 5



# The String data type: performance

---

String data type (in Java). Sequence of characters (immutable).

Java 7u5. Immutable char[] array, offset, length, hash cache.

Java 7u6. Immutable char[] array, hash cache.

operation	Java 7u5	Java 7u6
<b>length</b>	1	1
<b>indexing</b>	1	1
<b>concatenation</b>	$m + n$	$m + n$
<b>substring extraction</b>	1	$n$
<b>immutable?</b>	✓	✓
<b>memory</b>	$64 + 2n$	$56 + 2n$

# A Reddit exchange

---

I'm the author of the `substring()` change. As has been suggested in the analysis here there were two motivations for the change

- Reduce the size of `String` instances. Strings are typically 20-40% of common apps footprint.
- Avoid memory leakage caused by retained substrings holding the entire character array.



bondolo

Changing this function, in a bugfix release no less, was totally irresponsible. It broke backwards compatibility for numerous applications with errors that didn't even produce a message, just freezing and timeouts... All pain, no gain. Your work was not just vain, it was thoroughly destructive, even beyond its immediate effect.



cypherpunks

## Suffix sort

---


Q. How to efficiently form (and sort) suffixes in Java 7u6?

A. Define Suffix class à la Java 7u5 String representation.

```
public class Suffix implements Comparable<Suffix>
{
    private final String text;
    private final int offset;

    public Suffix(String text, int offset) {
        this.text = text;
        this.offset = offset;
    }

    public int length() { return text.length() - offset; }
    public char charAt(int i) { return text.charAt(offset + i); }
    public int compareTo(Suffix that) { /* see textbook */ }
}
```



## Suffix sort


---

Q. How to efficiently form (and sort) suffixes in Java 7u6?

A. Define Suffix class à la Java 7u5 String representation.

```
Suffix[] suffixes = new Suffix[n];
for (int i = 0; i < n; i++)
    suffixes[i] = new Suffix(s, i);

Arrays.sort(suffixes);
```



4<sup>th</sup> printing (2013)

Optimizations. [5× faster and 32× less memory than Java 7u5 version]


- Use 3-way string quicksort instead of Arrays.sort().
- Manipulate suffix offsets directly instead of via explicit Suffix objects.

# Suffix arrays: theory

**Conjecture.** [Knuth 1970] Impossible to compute suffix array in  $\Theta(n)$  time.

**Proposition.** [Weiner 1973] Can solve in  $\Theta(n)$  time (suffix trees).

**" has no practical virtue... but a historic  
monument in the area of string processing."**



## A Space-Economical Suffix Tree Construction Algorithm

EDWARD M. MCCREIGHT

*Xerox Palo Alto Research Center, Palo Alto, California*

**ABSTRACT.** A new algorithm is presented for constructing auxiliary digital search trees to aid in exact-match substring searching. This algorithm has the same asymptotic running time bound as previously published algorithms, but is more economical in space. Some implementation considerations are discussed, and new work on the modification of these search trees in response to incremental changes in the strings they index (the update problem) is presented.

## On-line construction of suffix trees<sup>1</sup>

Esko Ukkonen

Department of Computer Science, University of Helsinki,  
P. O. Box 26 (Teollisuuskatu 23), FIN-00014 University of Helsinki, Finland  
Tel.: +358-0-7084172, fax: +358-0-7084441  
Email: ukkonen@cs.Helsinki.FI

# Suffix arrays: practice

Applications. Bioinformatics, information retrieval, data compression, ...

Many ingenious algorithms.

- Constants and memory footprint very important.
- State-of-the art still changing.

year	algorithm	worst case	memory	
1991	<b>Manber–Myers</b>	$n \log n$	$8n$	see lecture videos
1999	<b>Larsson–Sadakane</b>	$n \log n$	$8n$	about 10x faster than Manber–Myers
2003	<b>Kärkkäinen–Sanders</b>	$n$	$13n$	
2003	<b>Ko–Aluru</b>	$n$	$10n$	
2008	<b>divsufsort2</b>	$n \log n$	$5n$	good choices (libdivsufsort)
2010	<b>sais</b>	$n$	$6n$	

# String sorting summary

# We can develop linear-time sorts.

- Key compares not necessary for string keys.
  - Use characters as index in an array.

We can develop sublinear-time sorts.

- Input size = total number of characters (not number of strings).
  - Not all of the characters have to be examined.

Long strings are rarely random in practice.

- Goal is often to learn the structure!
  - May need specialized algorithms.



© Copyright 2021 Robert Sedgewick and Kevin Wayne